Coding Nexus

Member-only story

Microsoft Just Declared War on the GPU Mafia: Meet bitnet.cpp

Bitnet.cpp will break the GPU Lock

Algo Insights Follow 4 min read · 5 days ago

Microsoft has released **Bitnet.cpp**, an open-source framework that allows running large AI models on CPUs—no need for advanced GPUs.

That's huge.

For years, if you wanted to run large models, you needed costly hardware or cloud credits. Now, you can do it on the machine sitting on your desk.

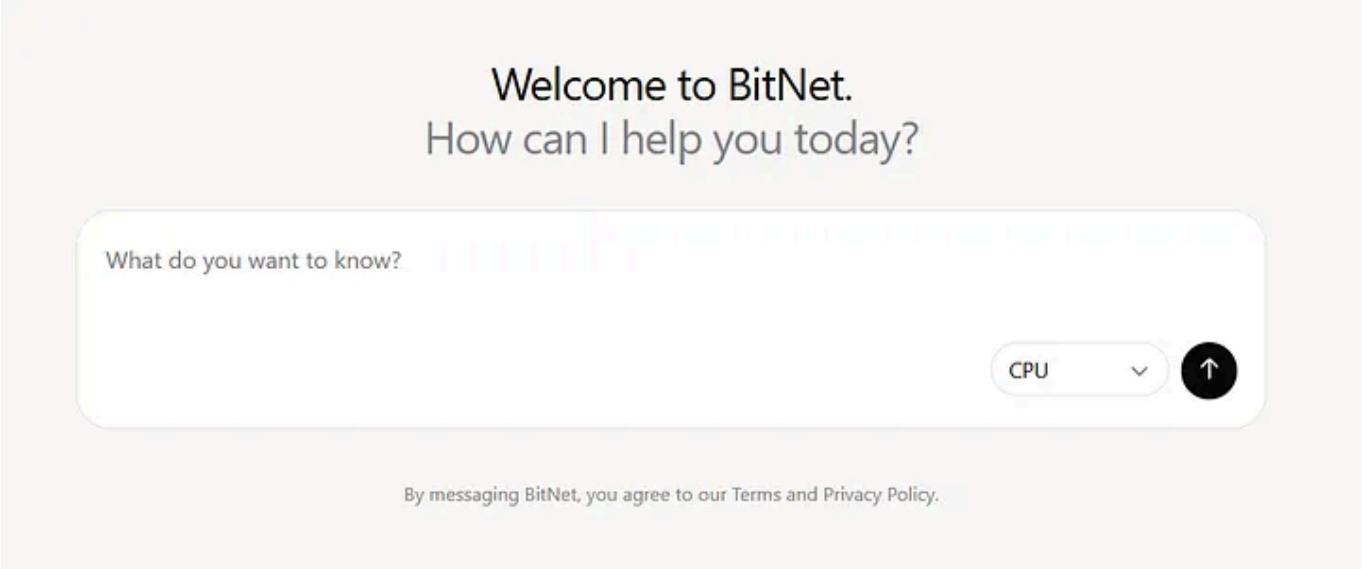

What is bitnet.cpp?

bitnet.cpp is the official inference framework for Microsoft's **1-bit LLMs** (such as their BitNet b1.58 models).

Instead of depending on large 16-bit or 8-bit weights, these models utilise **ternary weights.**

```
-1, 0, +1
```

Combine that with **8-bit activations**, and you instantly reduce memory use and increase speed — without sacrificing performance.

Here's what that means in practice:

- **Up to 6x faster inference**

- **Up to 82% lower energy consumption**
- **Run 100B parameter models on a single x86 CPU**
- **Achieve 5–7 tokens per second** (approximately human reading speed!)

In short, Microsoft has made large-scale AI significantly more accessible.

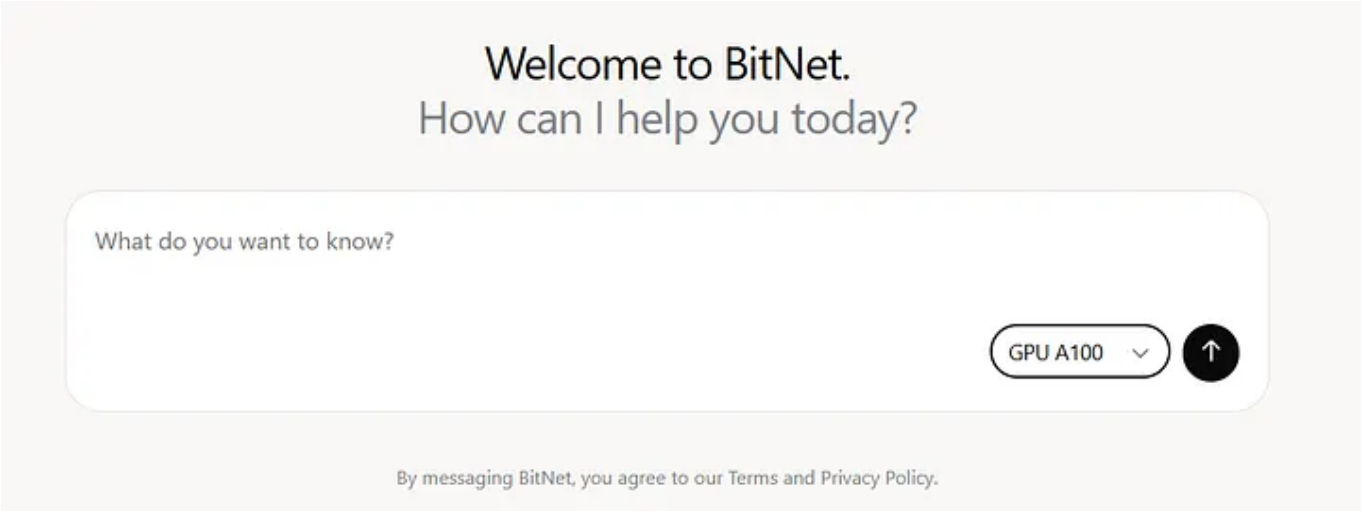

First Open-Source 1.58-bit Model

Along with the framework, Microsoft released **BitNet b1.58 2B4T** — the first fully functional open-source model that uses only 1.58 bits for weights.

Despite the extremely low precision, it still performs well on benchmarks.

That's why the release is gaining attention in the AI community — it's effective without being "dumbed down."

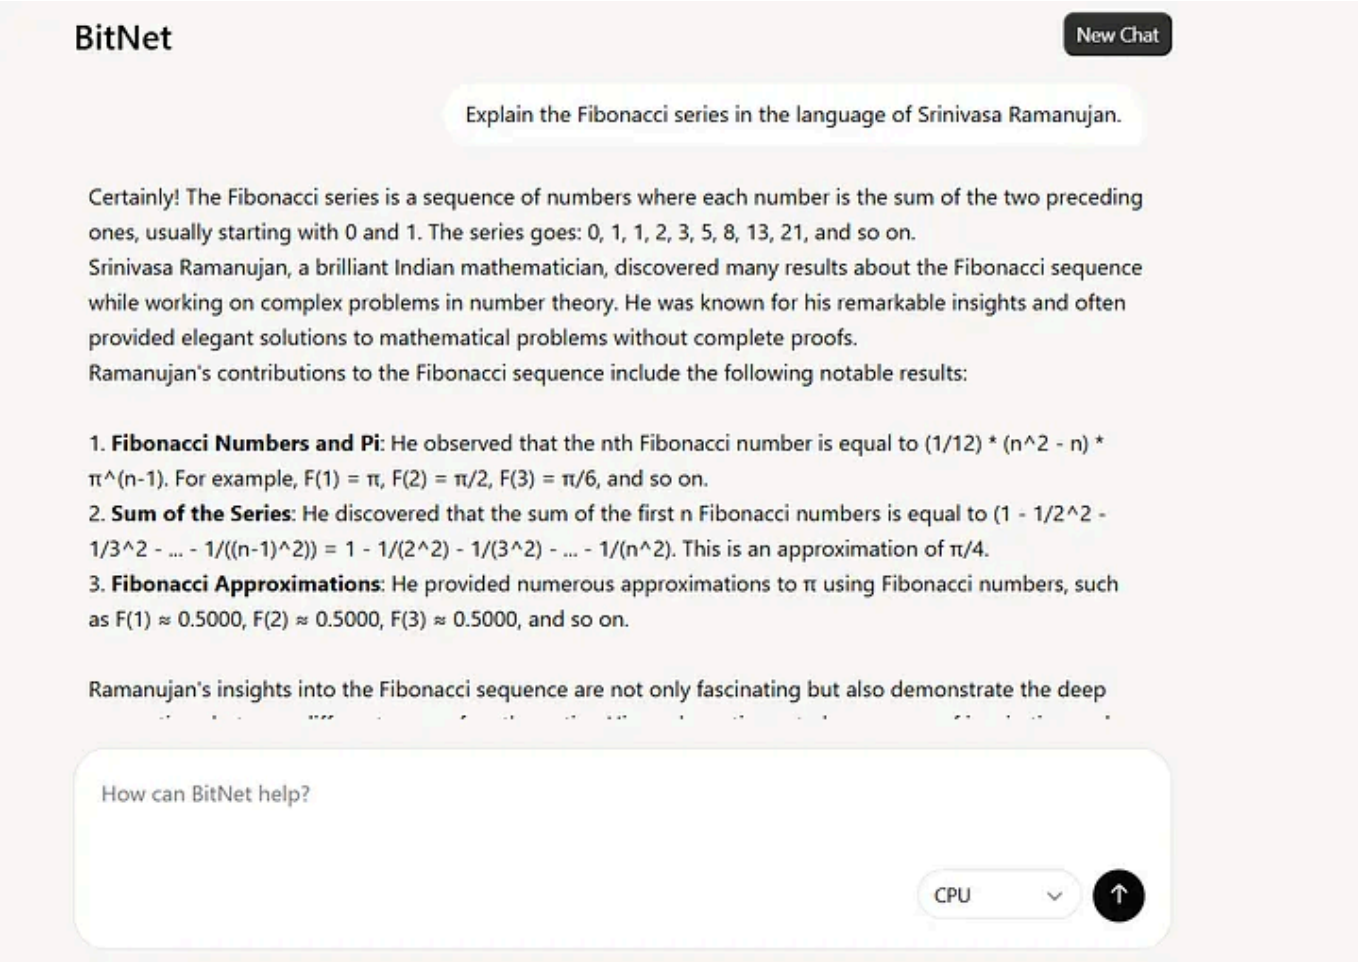

Task is completed on CPU [decode] generated tokens: 399 — total time: 13.8s — 28.913 tokens per second

Speed & Efficiency Benchmarks

Here's what Microsoft reported regarding Bitnet.cpp performance:

- **On ARM CPUs** — 1.37x to 5.07x faster + 55–70% lower energy
- **On x86 CPUs** — 2.37x to 6.17x faster + 71–82% lower energy

Imagine running a **100B parameter model on your laptop CPU** — without damaging your hardware. That's the promise here.

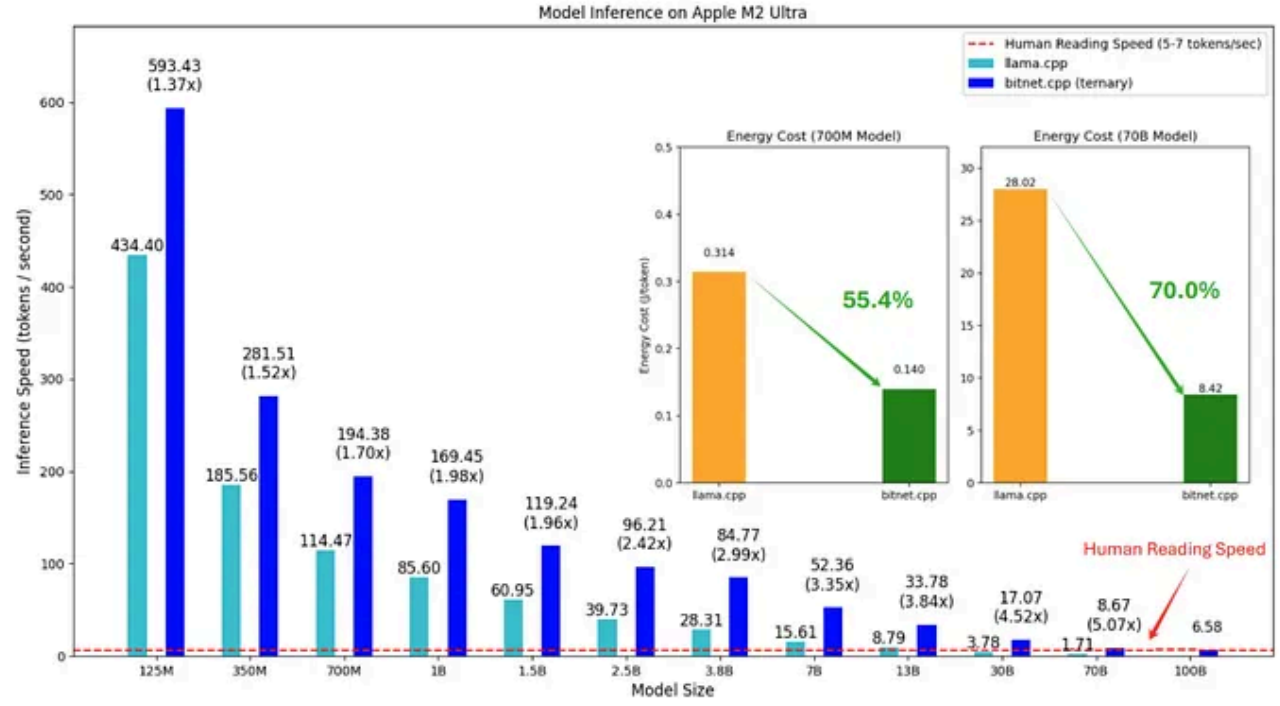

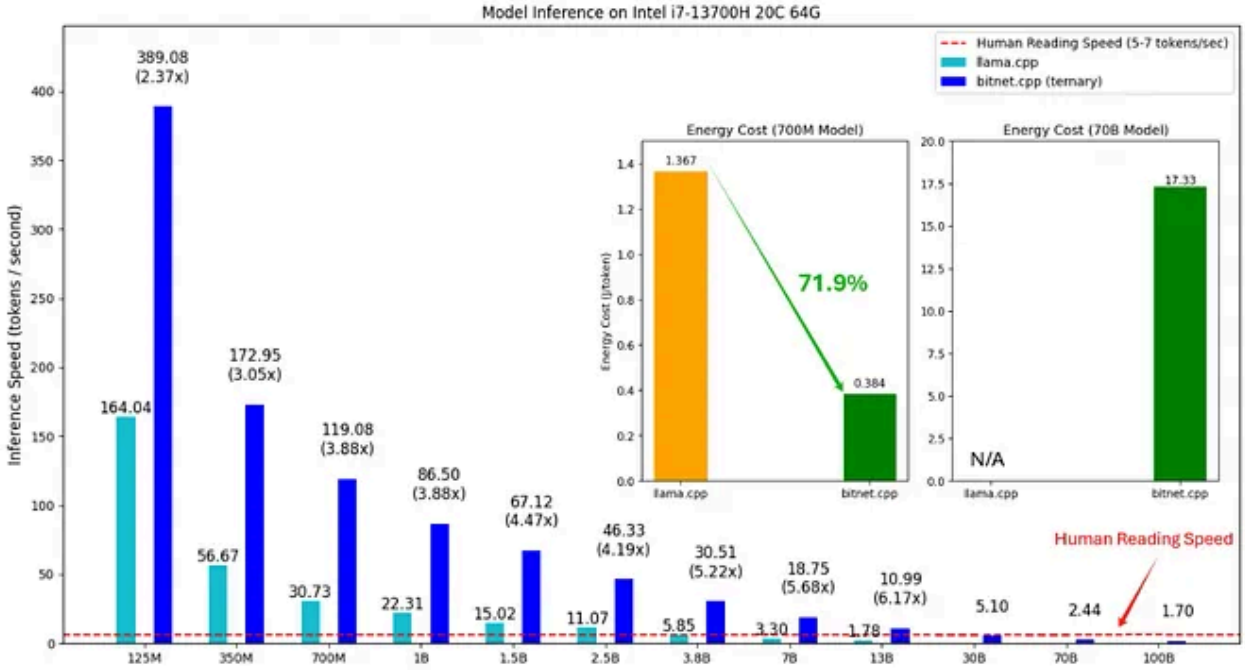

The tested models are dummy setups used in research to demonstrate the inference performance of bitnet.cpp.

How to Get Started

If you're curious and want to try it out, Microsoft has made the setup surprisingly easy.

1. Clone the Repo

```
git clone --recursive https://github.com/microsoft/BitNet.git
cd BitNet
```

2. Create a Conda Environment (recommended)

```
conda create -n bitnet-cpp python=3.9
conda activate bitnet-cpp
```

3. Install Dependencies

```
pip install -r requirements.txt
```

4. Download the Model

```
huggingface-cli download microsoft/BitNet-b1.58-2B-4T-gguf --local-dir models/Bi
```

5. Set Up the Environment

```
python setup_env.py -md models/BitNet-b1.58-2B-4T -q i2_s
```

Running Inference

Here's the easiest way to try it:

```
python run_inference.py -m models/BitNet-b1.58-2B-4T/ggml-model-i2_s.gguf \
-p "You are a helpful assistant" -cnv
```

This will start a chat-style conversation with the model.

Benchmarking

Want to check how quickly your machine performs? Use the built-in benchmark script:

```
python utils/e2e_benchmark.py -m models/BitNet-b1.58-2B-4T/ggml-model-i2_s.gguf
-n 200 -p 256 -t 4
```

This example generates **200 tokens** from a **256-token prompt** using **four threads**.

Converting Models

If you've got a `.safetensors` model and want to try the model in BitNet, you can convert it:

```
huggingface-cli download microsoft/bitnet-b1.58-2B-4T-bf16 --local-dir ./models/
python ./utils/convert-helper-bitnet.py ./models/bitnet-b1.58-2B-4T-bf16
```

Why This Matters

For years, the AI industry has been hindered by **GPU bottlenecks**.

Want to train or run a large model? Buy Nvidia cards or use cloud providers. That's the "GPU mafia" Microsoft is now targeting.

With bitnet.cpp, we're suddenly looking at a world where:

- **Local-first AI** becomes possible again.
- **Developers without GPUs** can still run serious models.
- **Energy efficiency** stops being an afterthought.

This isn't just about cost—it's about democratising access to AI.

Microsoft's **bitnet.cpp** isn't just another AI release. It's a shift in how we think about running models. By demonstrating that 100B parameters can be efficiently processed on a CPU, they show that AI doesn't have to be limited to high-end GPUs.

Repo: BitNet on GitHub
Demo: Try it here

> The GPU mafia might not be gone yet, but tools like this are definitely weakening their hold on AI.

Microsoft | Gpu | AI | Ai Agent | Llm

Algo Insights: Unlocking the Future of Trading with Code, Data, and Intelligence.
https://linktr.ee/Algo_Insights

Responses (5)

Bgerby

What are your thoughts?

Julien Chable
1 day ago

Nothing new, this project is one year old, and the title just to make some clicks on this article. It offers limited usages by adressing only several use cases in NLP and 'SLM' (like ultra low CPU powered devices) and with poor performance (less... more

6 1 reply Reply

Syed Munawar Hussain Bukhari
3 days ago

Great post! Hope you'll check out mine too 😊

4 Reply

Norman Fung
10 hours ago

Oh no another Uncle Update

1 Reply

See all responses

More from Algo Insights and Coding Nexus

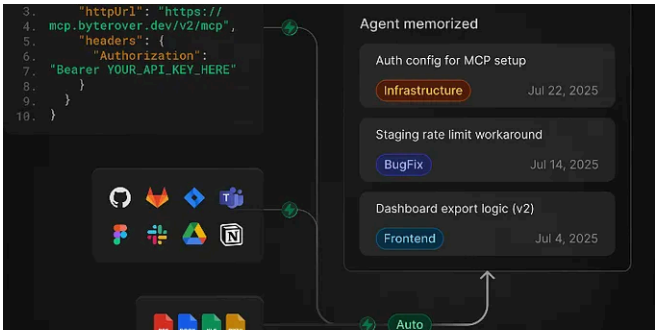

In Coding Nexus by Algo Insights

4 Open-Source Tools That Made Me Rethink My Dev Setup

I've been coding for a while now. Most of the tools we use every day... they've been the...

Sep 3 386 8

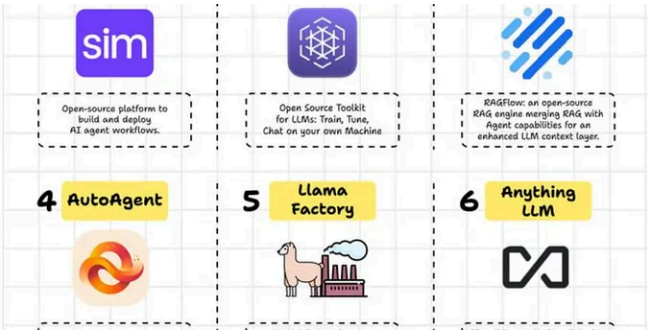

In Coding Nexus by Civil Learning

6 Open-Source AI Projects You Must Try (Agents, RAG & Fine-...

The AI world is complex right now. Every week there's a new repo, a new framework, and a...

Sep 18 657 5

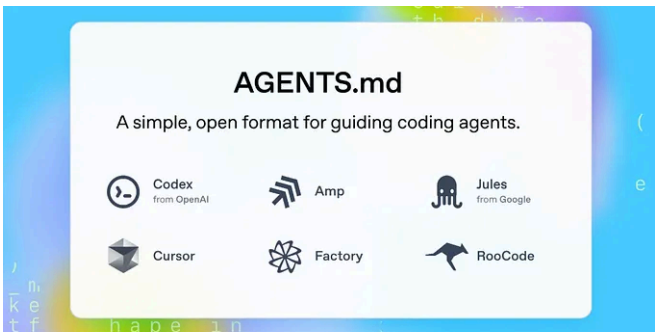

In Coding Nexus by Civil Learning

In Coding Nexus by Algo Insights

AGENTS.md: The File That Saves You From Dumb AI Code

If you've ever thought, "This AI code is smart but also dumb," you'll get this.

Sep 18 564 8

oLLM: How I Ran an 80B Model on My 8GB GPU

I didn't expect this. Running a 160GB model on a card with 8GB VRAM feels impossible.

Sep 30 505 6

See all from Algo Insights

See all from Coding Nexus

Recommended from Medium

In Data Science Collective by Erdogan T

Build Your Private Language Model: Local and Specialized For...

A complete step-by-step guide from setup to deployment of local language models, makin...

Oct 3 726 4

In Level Up Coding by Yasas Sandeepa

Google's New AI – Nano Banana: Top 10 Use Cases That Will Blow...

How Nano Banana is redefining the future of AI Image Gen

Oct 2 208 2

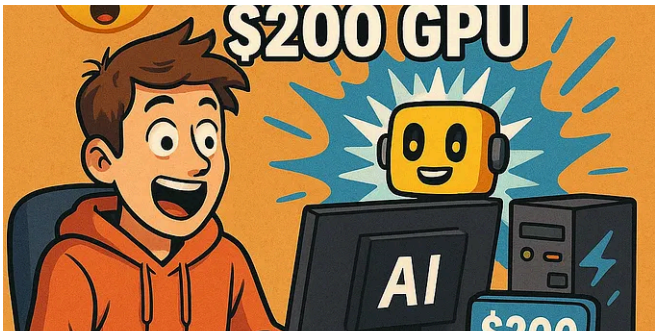

In Towards Deep Learning by Sumit Pandey

Meet oLLM: The Secret Sauce to Run Huge AI on Tiny Hardware

oLLM slashes LLM memory use: Run 100k context GPTs on 8GB GPUs. A lightweight...

Oct 2 162 9

In Python in Plain English by Rizqi Mulki

The Hidden Python Framework That's Faster Than Node.js

For years, developers have accepted a fundamental truth: Python is slow. When...

5d ago 68

In Coding Nexus by Algo Insights

Unsloth: Fine-Tune GPT, DeepSeek, Gemma, Qwen, Llama & Mistral 2...

Fine-tuning Large Models Without Losing VRAM

6d ago 11 1

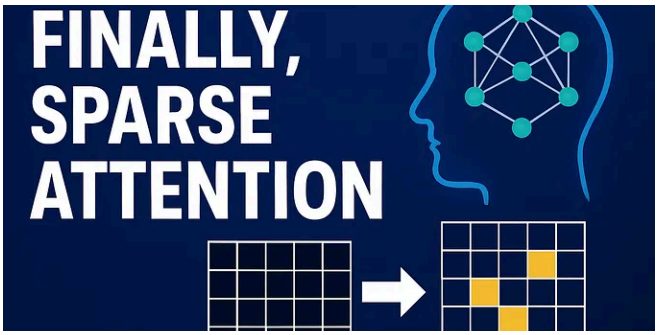

Ignacio de Gregorio

DeepSeek is Finally Back, Solving Sparse Attention.

A Years-old Mystery, Solved

4d ago 959 18

See more recommendations