

Data, Analytics & A... · [Follow publication](#)

# A Journey from AI to LLMs and MCP — 10 — Sampling and Prompts in MCP — Making Agent Workflows Smarter and Safer

4 min read · May 20, 2025



Alex Merced



[Follow](#)

Listen

Share

More



## Free Resources

- [Free Apache Iceberg Course](#)
- [Free Copy of “Apache Iceberg: The Definitive Guide”](#)
- [2025 Apache Iceberg Architecture Guide](#)

- [How to Join the Iceberg Community](#)
- [Iceberg Lakehouse Engineering Video Playlist](#)
- [Ultimate Apache Iceberg Resource Guide](#)

We've now seen how the Model Context Protocol (MCP) allows LLMs to read resources and call tools — giving them access to both data and action.

But what if your MCP server needs the LLM to make a decision?

What if it needs to:

- Analyze a file before running a tool?
- Draft a message for approval?
- Ask the model to choose between options?

That's where Sampling comes in.

And what if you want to give the user — or the LLM — reusable, structured prompt templates for common workflows?

That's where Prompts come in.

In this final post of the series, we'll explore:

- How sampling allows servers to request completions from LLMs
- How prompts enable reusable, guided AI interactions
- Best practices for both features
- Real-world use cases that combine everything we've covered so far

## **What Is Sampling in MCP?**

Sampling is the ability for an MCP server to ask the host to run an LLM completion — on behalf of a tool, prompt, or workflow.

It lets your server say:

*“Hey, LLM, here’s a prompt and some context. Please respond.”*

## Why is this useful?

- You can generate intermediate reasoning steps
- Let the model propose actions before executing them
- Create more natural multi-turn agent workflows
- Maintain human-in-the-loop approval and visibility

## Sampling Flow

Here's the typical lifecycle:

1. The server sends a `sampling/createMessage` request
2. The host (Claude Desktop, etc.) can review or modify the prompt
3. The host runs the LLM completion
4. The result is sent back to the server

*This architecture puts control and visibility in the hands of the user, even when the agent logic runs server-side.*

## Message Format

Here's an example `sampling/createMessage` request:

```
{  
  "messages": [  
    {  
      "role": "user",  
      "content": {  
        "type": "text",  
        "text": "Please summarize this log file."  
      }  
    }  
  ],  
  "systemPrompt": "You are a helpful developer assistant.",  
  "includeContext": "thisServer",  
  "maxTokens": 300  
}
```

The host chooses which model to use, what context to include, and whether to show the prompt to the user for confirmation.

Response:

```
{  
  "model": "claude-3-sonnet",  
  "role": "assistant",  
  "content": {  
    "type": "text",  
    "text": "The log file contains several timeout errors and warnings related  
  }  
}
```

Now the server can act on that response — log it, return it as tool output, or chain it into another step.

## Best Practices for Sampling

### Best Practice Why It Matters

- Use clear system prompts Guides model behavior contextually
- Limit tokens Prevent runaway completions
- Structure responses Enables downstream parsing (e.g. JSON, bullets)
- Include only relevant context Keep prompts focused and cost-effective
- Respect user control The host mediates the actual LLM call

## What Are Prompts in MCP?

Prompts are reusable, structured templates that servers can expose to clients.

Think of them like slash commands or predefined workflows:

- Pre-filled with helpful defaults
- Accept arguments (e.g. “project name”, “file path”)
- Optionally include embedded resources
- Surface in the client UI

Prompts help users and LLMs collaborate efficiently by standardizing useful tasks.

## Prompt Structure

Prompts have:

- A name (identifier)
- A description (for discovery)
- A list of arguments (optional)
- A template for generating messages

Example:

```
{  
  "name": "explain-code",  
  "description": "Explain how this code works",  
  "arguments": [  
    {  
      "name": "language",  
      "description": "Programming language",  
      "required": true  
    },  
    {  
      "name": "code",  
      "description": "The code to analyze",  
      "required": true  
    }  
  ]  
}
```

Clients use:

- `prompts/list` to discover prompts
- `prompts/get` to resolve a prompt and arguments into messages

## Dynamic Prompt Example

A server might expose:

```
{  
  "name": "analyze-logs",  
  "description": "Summarize recent logs and detect anomalies",  
  "arguments": [  
    {  
      "name": "log_type",  
      "description": "Type of log to analyze",  
      "required": true  
    },  
    {  
      "name": "start_time",  
      "description": "Start time of the log analysis",  
      "required": false  
    },  
    {  
      "name": "end_time",  
      "description": "End time of the log analysis",  
      "required": false  
    }  
  ]  
}
```

```
        "name": "timeframe",
        "required": true
    }
]
}
```

When the user (or LLM) runs it with:

```
{
  "timeframe": "1h"
}
```

The resolved prompt could include:

- A message like: “Please summarize the following logs from the past hour.”
- An embedded resource (e.g. `logs://recent?timeframe=1h`)
- Output ready for sampling

## **Sampling + Prompts = Dynamic Workflows**

Combining prompts + sampling + tools unlocks real agent behavior.

Example Workflow:

- **User selects prompt:** “Analyze logs and suggest next steps”
- Server resolves the prompt and calls sampling/createMessage
- **LLM returns:** “The logs show repeated auth failures. Suggest checking OAuth config.”
- Server calls tools/call to run `check_auth_config`
- LLM reviews the result and writes a summary

All controlled via:

- Standardized MCP messages
- User-visible approvals

- Modular server logic

## Security and Control

## Why These Matter for AI Agents

## Wrapping It All Together

Over this 10-part series, we've explored the full landscape of AI agent development using MCP:

- LLMs and how they work
- Fine-tuning, prompting, and RAG
- Agent frameworks and limitations

- MCP's architecture and interoperability
- Resources and tools
- Prompts and sampling

MCP gives us standardized, modular building blocks for creating AI agents that are:

- Portable across environments
- Decoupled from model providers
- Secure, observable, and controlled

AI

Model Context Protocol

Machine Learning

Follow

## Published in Data, Analytics & AI with Dremio

375 followers · Last published Oct 14, 2025

The Intelligent Lakehouse Platform

Follow

## Written by Alex Merced

633 followers · 14 following

I'm a tech, development and data enthusiast who has a lot to say. You can find all my blogs, videos and podcasts at [AlexMerced.com](http://AlexMerced.com)

## No responses yet



Bgerby

What are your thoughts?

## More from Alex Merced and Data, Analytics & AI with Dremio

 In Data, Analytics & AI with Dremio by Alex Merced 

### Building a Basic MCP Server with Python

Free Resources

Apr 4     163     18



...

 In Data, Analytics & AI with Dremio by Alex Merced 

## The State of Apache Iceberg v4—October 2025 Edition

Get Data Lakehouse Books:

Oct 14  1



...

 In Data, Analytics & AI with Dremio by Alex Merced 

## 2025 Guide to Architecting an Iceberg Lakehouse

Blog: What is a Data Lakehouse and a Table Format?

Dec 10, 2024  96



...

 Alex Merced 

## The 2025 & 2026 Ultimate Guide to the Data Lakehouse and the Data Lakehouse Ecosystem

Join the Data Lakehouse Community

Sep 23  5



...

[See all from Alex Merced](#)

[See all from Data, Analytics & AI with Dremio](#)

Recommended from Medium

 Riccardo Tartaglia

## 5 Essential MCP Servers Every Developer Should Know

I've been experimenting with Model Context Protocol servers for a few months now, and I have to say, they've changed the way I work.

Oct 11  22  3



...

 Vandana Maurya

## Confusion Matrix—Complete Beginner to Advanced Guide

Here, I am sharing easy-to-read, beginner-friendly, complete note on Confusion Matrix you can go through just before your interview or...

Apr 29



...

In Towards AI by Ahmed Boulahia

## Best Open-Source Embedding Models for RAG

High-Performance Open-Source Embedding Models for RAG Pipelines, Multilingual NLP, and Arabic Text

Oct 14 201



...

In Coding Nexus by Civil Learning

## MarkItDown: Convert Anything into Markdown—the Smart Way to Feed LLMs

You know that feeling when you're trying to feed a PDF or a Word document into an LLM, and it just doesn't understand what's going on...

⭐ 6d ago ⌚ 125 💬 1



...

Matteo Ferruccio Andreoni

## How to Fine-Tune LLMs Locally: The Complete LoRA Guide

Master LoRA Fine-Tuning on Apple Silicon with minimal memory usage. From setup to deployment in one comprehensive tutorial.

3d ago ⌚ 87



...



Alex Suzuki

## Cretaceous Software Engineering—Why I don't use AI to write code.

👋 Hi! I'm a Cretaceous Software Engineer.

5d ago    641    28



...

[See more recommendations](#)