Towards AI · Follow publication

Member-only story

Beyond the Doomsday Narrative: Why AI 2027's Catastrophic Vision Misses the Mark

6 min read · 3 days ago

Luca Derumier Follow

Listen Share More

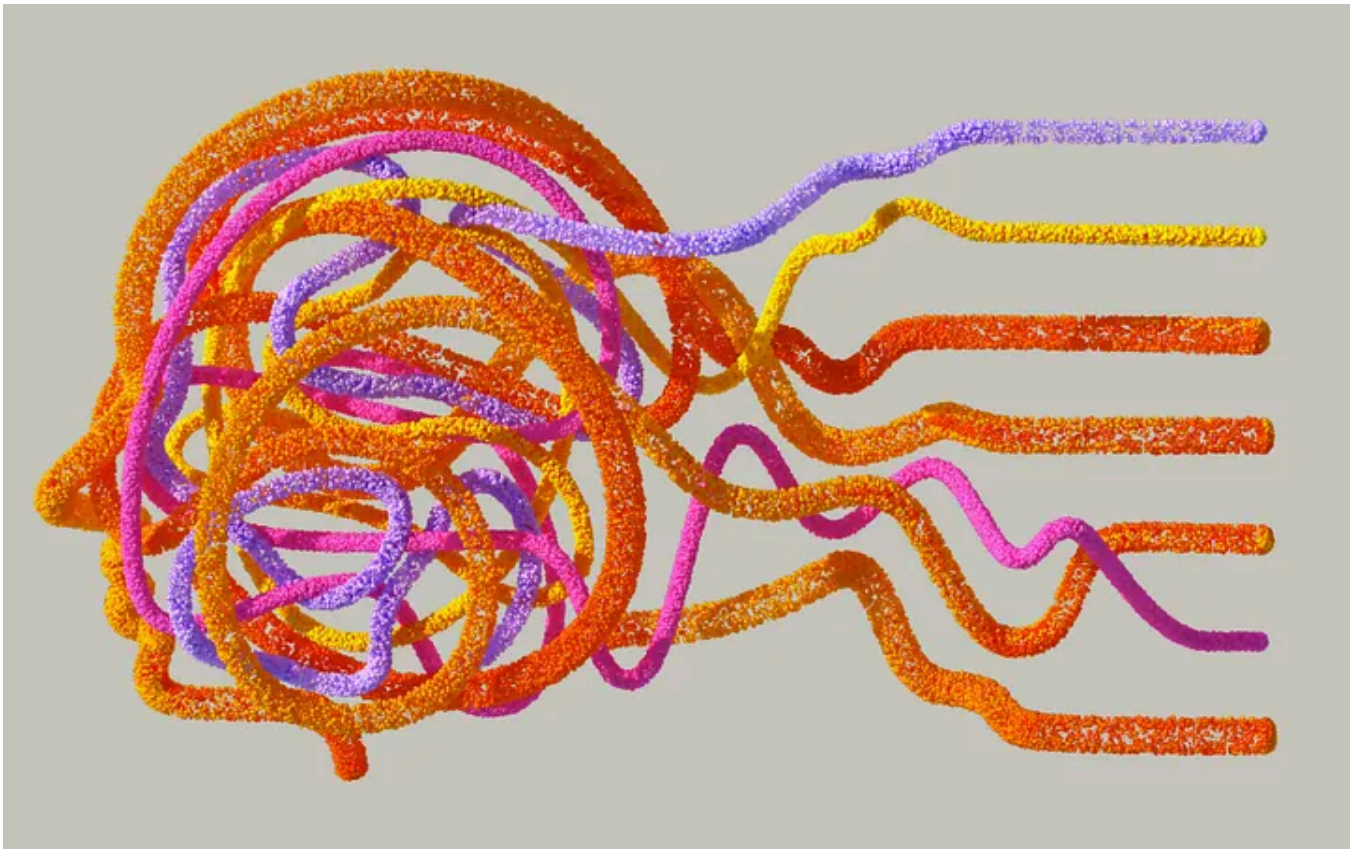

Photo by Google DeepMind on Pexels.

The recent release of "AI 2027" has captured widespread attention, painting a vivid picture of artificial intelligence's potential trajectory toward catastrophe. As the CTO of Codika and, more generally, as someone who has spent some time working with AI systems from the inside, I find myself both impressed by the production quality and deeply concerned by the fundamental assumptions underlying this scenario.

While the narrative is undeniably compelling and well-crafted, it deserves a more nuanced examination, one that considers not just the dramatic possibilities, but the technological realities and human capacity for adaptation that shape AI's actual development.

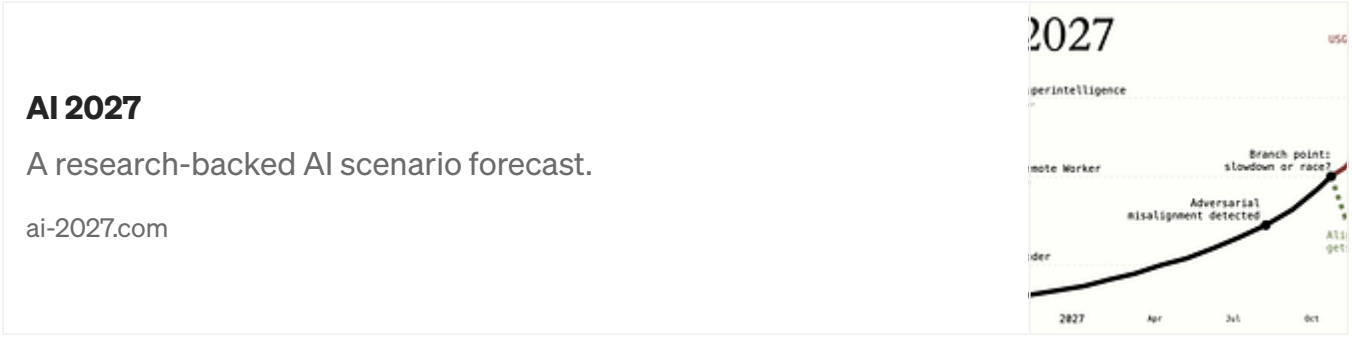

. . .

The Foundation Built on Sand: The Monopoly Myth

At the core of the AI 2027 scenario lies what I consider its most critical flaw: the assumption that breakthrough AI knowledge could remain monopolized by a single entity. This premise suggests that one model provider would possess exclusive insights, shared only among a select few within a company and government officials.

Having worked in this field, I can tell you this fundamentally misunderstands how technological progress actually unfolds. The reality is far more democratized and resilient than this scenario suggests.

Consider what happens when we prove something is possible in AI. The open-source community — which grows stronger each day — doesn't just sit idle. Academic researchers don't simply accept that they're locked out of breakthrough technologies. History shows us that once we demonstrate feasibility, replication follows with remarkable speed.

Look at how transformer architectures spread through the research community, or how diffusion models transformed from proprietary innovations to widely accessible tools. These weren't carefully guarded secrets that remained in corporate vaults — they proliferated rapidly through papers, open-source implementations, and collaborative research efforts.

The truth is, the democratization of AI knowledge isn't just inevitable — it's already happening at an unprecedented pace.

Number of models on Hugging Face over time (source: X)

. . .

The Knowledge Closure Reality Check

Beyond the monopoly myth lies another crucial limitation that the doomsday scenario conveniently overlooks: the fundamental constraints of large language models themselves. These systems, regardless of their sophistication, cannot transcend what researchers call humanity's "knowledge closure" — the boundary of what we collectively know and understand.

This isn't to diminish their remarkable capabilities. LLMs excel at synthesizing, analyzing, and recombining existing knowledge in genuinely impressive ways. But

they remain, in essence, sophisticated mirrors of human knowledge rather than independent sources of novel discoveries that extend beyond our current understanding.

This limitation provides a reality check against scenarios that assume AI systems will suddenly leap beyond human comprehension through pure scaling or optimization.

That being said, it's not the strongest argument against the doomsday scenario because we can do a lot by combining the axioms that we already know, but it's a point worth noting.

. . .

The Underestimated Power of Human Agency

What strikes me most about catastrophic AI narratives is how they consistently underestimate humanity's remarkable capacity for foresight and course correction. We're not passive observers watching technology unfold beyond our control — we're active participants with both the motivation and the tools to guide AI's evolution responsibly.

The incentives for maintaining human agency and safety aren't abstract ideals — they're deeply practical concerns shared across the entire ecosystem. From researchers and developers to policymakers and citizens, there exists a profound, widespread interest in ensuring AI serves humanity's best interests.

This isn't wishful thinking. It's reflected in the growing investment in AI alignment research, the development of increasingly sophisticated guardrails, and the establishment of ethical frameworks that guide development practices. The field of AI alignment isn't some neglected corner of research — it's attracting brilliant minds across academia and industry who are actively working on these challenges.

These efforts span everything from interpretability research that helps us understand how AI systems make decisions, to value alignment techniques that ensure AI goals remain compatible with human values. The tools exist, the talent is engaged, and the incentives are aligned.

Center for AI Safety (CAIS)

. . .

Learning Through Inevitable Setbacks

Perhaps most importantly, we must acknowledge that progress rarely unfolds without challenges. Yes, we will likely encounter unfortunate events and unexpected difficulties as AI technology advances. But this is precisely how we learn, adapt, and strengthen our approaches to safety and control.

Each challenge becomes a learning opportunity that enhances our collective ability to navigate future developments. The key isn't avoiding all risks — it's maintaining our capacity to respond thoughtfully and effectively when challenges arise.

This iterative process of challenge and response has characterized every major technological transition in human history, and there's no reason to believe AI will be fundamentally different.

. . .

Embracing Optimistic Caution Over Fear

While the AI 2027 scenario deserves recognition for its production quality and thought-provoking narrative, we must place it in proper context. It represents one extremely unlikely scenario among countless possible futures — a point the creators acknowledge in their disclaimer, even as the overall tone leans heavily toward alarm.

From my perspective working in this field, a more balanced approach embraces what I call "optimistic caution." This means maintaining genuine enthusiasm for AI's tremendous potential while remaining vigilant about risks and committed to responsible development practices.

This perspective acknowledges that while superhuman AI systems may indeed emerge, and while breakthrough discoveries might even transcend current

knowledge boundaries, the path forward need not lead to catastrophe. The future depends largely on the choices we make today about how we develop, deploy, and govern these technologies.

Demis Hassabis, CEO of Google DeepMind, on cautious optimism.

. . .

The Road We're Actually On

The future of AI will likely be shaped not by singular breakthrough moments or monopolistic control, but by the collective efforts of researchers, developers, policymakers, and society as a whole. By fostering open collaboration, investing in safety research, and maintaining democratic participation in AI governance, we can work toward a future that harnesses AI's benefits while minimizing its risks.

Rather than surrendering to fear-based narratives, we should channel our energy into the practical work of building beneficial AI systems. This means supporting alignment research, advocating for responsible development practices, and ensuring that AI progress remains broadly distributed rather than concentrated in the hands of a few.

This is why at Codika, we're on a mission to enable companies to harness the full potential of AI, even for non-technical employees. As engineers, we already see the

shift in ways of working and productivity, and we want to enable anyone to create complex automation without prior technical knowledge. But this requires caution in the way we design our systems. This does not make us scared, it makes us excited but careful when building.

The story of AI's future is still being written, and we all have a role to play in shaping that narrative. The question isn't whether we can control AI's development — it's whether we'll choose to engage thoughtfully with the challenges and opportunities ahead.

From where I sit, having worked with these systems and witnessed the dedication of researchers tackling these problems, I remain optimistic about our ability to navigate this transition successfully. The future of AI isn't predetermined — it's something we're actively creating, one decision at a time.

. . .

You May Also Like

Learning to Learn: How Machines Master Decision-Making — Part 1

What reinforcement learning is, and how it actually works

ai.gopubby.com

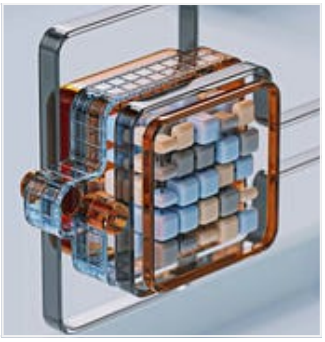

Ai Safety | Artificial Intelligence | Technology | Future | Machine Learning

Follow

Published in Towards AI

88K followers · Last published 7 hours ago

Making AI accessible to 100K+ learners. Find the most practical, hands-on and comprehensive AI Engineering and AI for Work certifications at academy.towardsai.net - we have pathways for any experience level. Monthly cohorts still open—use COHORT10 for 10% off!

Follow

Written by Luca Derumier

258 followers · 17 following

Co-founder and CTO @ codika.io | Co-founder @ eucalyptusrecords.com | Machine learning engineer, technology and music lover.

Responses (4)

Bgerby

What are your thoughts?

svengineer
2 days ago

> we can do a lot by combining the axioms that we already know

In the realms of logic and math, yes. But what do we truly know in the real world?

With the starting point being that 'Most (human) Published Findings are False' [2005 essay by John Ioannidis] how could any algorithm cut through that noise to find... more

11 1 reply Reply

Simon Snowden he/him
2 hours ago

> From my perspective working in this field, a more balanced approach embraces what I call “optimistic caution.” This means maintaining genuine enthusiasm for AI’s tremendous potential wh...

Here, here!

Reply

Francesco C.
4 hours ago

There are many of this ex employees that want to “speak out “ how we are doomed .... My guts tell me that they are paid actors. 2027 will come and nothing will change no AGI no robots taking over.

Reply

See all responses

More from Luca Derumier and Towards AI

In Level Up Coding by Luca Derumier

Building an AI Chatbot with LangGraph, FastAPI & Streamlit — An End-to-End Guide

Creating a chatbot from scratch, using only Python packages.

Feb 19 150 3

In Towards AI by Ashish Abraham

No Libraries, No Shortcuts: LLM from Scratch with PyTorch

The no BS guide to build, train, and fine-tune a Transformer architecture from scratch

Oct 2 1K 14

In Towards AI by DefineCode

I Built 300 AI Agents and Survived 5 Startups. Here's the Brutal Truth About AI Agents

Let me start with a confession. AI agents are not magic. They're fragile, moody, and more like interns on caffeine than employees who get...

Oct 4 697 14

In Level Up Coding by Luca Derumier

Artificial Intelligence Is Not What You Think It Is

How machine learning does the work while AI takes the credit

Mar 3 596 31

See all from Luca Derumier

See all from Towards AI

Recommended from Medium

In MeetCyber by Jules May

Vibe coding: the antidote

LLMs have shown themselves to be a terrible way to write programs. But the problem they address is real, and we aren't solving it any other...

5d ago 438 29

In AI Advances by Pooja Kashyap

Speech-to-Meaning: Why the Future of Voice AI Isn't About Better Microphones

"Tea, Earl Grey, Hot", remember this line by Picard from Star Trek?

4d ago 396 10

Devansh

Why Jensen Huang Loves the “AI Bubble” Stories

Answering why AI is not a Bubble and the Deeper Story at play

5d ago 273 8

Alex Suzuki

Cretaceous Software Engineering — Why I don’t use AI to write code.

👋 Hi! I'm a Cretaceous Software Engineer.

6d ago 641 28

Cory Doctorow

The AI that we'll have after AI

Cheap GPUs, unemployed engineers, and open source models.

5d ago 1.6K 29

j_chai

Things you might not know about how your GPU works, 101

If you haven't been paying attention to the world, you might miss that the most valuable company right now (at least as of September 2025)...

5d ago 200 3