

★ Member-only story

Timeplus Proton: The Lightning-Fast SQL Engine for Real-Time Streams

4 min read · 2 days ago



Algo Insights

Following

Listen

Share

More

Real-time data is a mess. If you've ever tried building a streaming pipeline, you know the chaos. Kafka for streams, Flink or ksqlDB for processing, ClickHouse for analytics, connectors for... everything else.

Before long, you're knee-deep in YAML and microservices, wondering if this was meant to be *fun*.

Then I discovered **Timeplus Proton**. It's essentially: "*What if stream processing was just SQL — and it ran like a rocket?*"



Fastest SQL pipeline engine for stream processing, analytics, observability and AI

• • •

So, What's Proton?

Proton is a single binary — written in C++ — that performs stream processing, analytics, observability, and even a touch of AI magic.

No JVM.

No ZooKeeper.

No massive cluster setup.

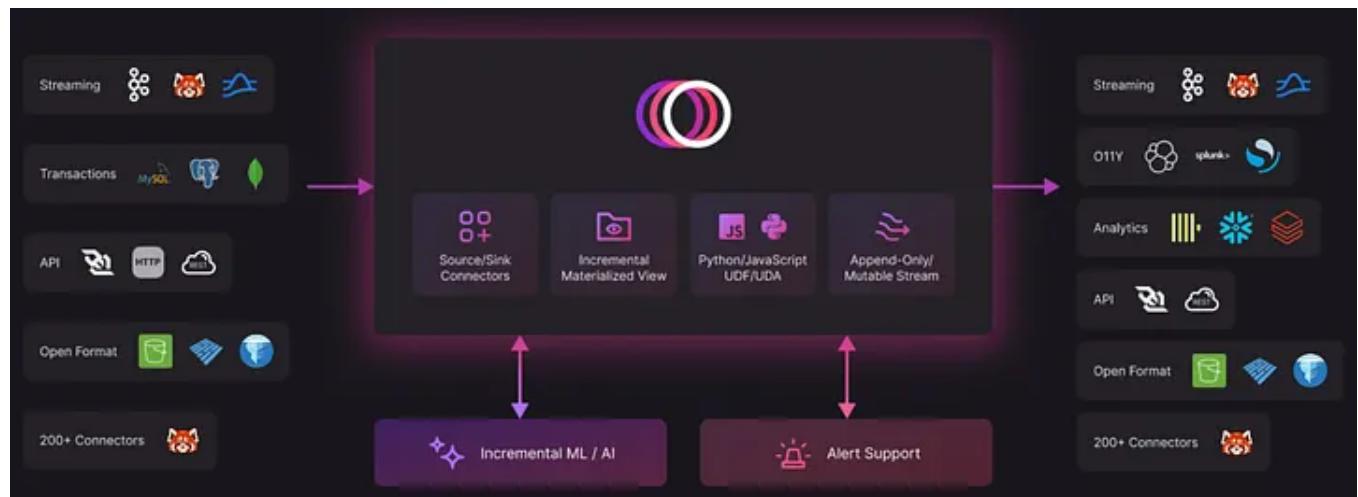
You just run:

```
curl https://install.timeplus.com/oss | sh
```

And you're live. Literally.

Behind the scenes, it uses the **ClickHouse engine**, which is already one of the fastest databases available. However, Proton adds all the streaming features that ClickHouse doesn't offer — joins, windows, materialised views, watermarks, and more.

It's like ClickHouse has learned to listen to data as it occurs.



The Cool Part — It's Actually Fast

I'm not talking “marketing fast.”

I mean, “90 million events per second on a MacBook” fast.

Proton's written in C++ and uses SIMD optimisation (which basically means it squeezes more work out of your CPU).

And because it's a single binary under 500MB, it doesn't chew through memory or spin up ten background services to say "hello."

You can even run it on an AWS t2.nano — that's one vCPU and 0.5GB memory. Try doing that with Flink.

• • •

ClickHouse vs Proton (in simple terms)

Here's how I see it:

- **ClickHouse** is for *big historical analytics*. It's a data warehouse.
- **Proton** is for *real-time decisions*. It's a stream engine.

ClickHouse can read from Kafka, sure, but it doesn't *think in streams*. It thinks in tables. Proton lives in both worlds — it unifies real-time and historical data. You can detect something now and query it later, without juggling two systems.

It's similar to ClickHouse and Flink merging into a small binary.

• • •

Getting Hands-On

Once installed, open Proton's SQL client and begin writing SQL—no new language to learn.

Example — reading from Kafka and writing to ClickHouse:

```
CREATE EXTERNAL STREAM aws_msk_stream (
    device string,
    temperature float
)
SETTINGS
    type='kafka',
    brokers='prefix.kafka.us-west-2.amazonaws.com:9098',
    topic='topic',
```

```
security_protocol='SASL_SSL',
sasl_mechanism='AWS_MSK_IAM';
```

Then send the results to ClickHouse:

```
CREATE EXTERNAL TABLE ch_aiven
SETTINGS type='clickhouse',
address='abc.aivencloud.com:28851',
user='avnadmin',
password='..',
secure=true,
table='events';
```

And finally, create a view that aggregates data every 10 seconds:

```
CREATE MATERIALIZED VIEW mv_msk2ch INTO ch_aiven AS
SELECT window_start AS timestamp, device, avg(temperature) AS avg_temperature
FROM tumble(aws_msk_stream, 10s)
GROUP BY window_start, device;
```

That's real-time ETL – no frameworks, no YAML, just SQL.

• • •

Don't Have Kafka? No Problem.

You can generate random data for testing.

```
CREATE RANDOM STREAM devices(
device string default 'device'||to_string(rand()%4),
temperature float default rand()%1000/10);

SELECT device, count(*), min(temperature), max(temperature)
FROM devices GROUP BY device;
```

You'll see a live stream of fake IoT data passing by. It's oddly satisfying.

device	count()	min(temperature)	max(temperature)
device0	2256	0	99.6
device1	2260	0.1	99.7
device3	2259	0.3	99.9
device2	2225	0.2	99.8

Feels like magic.

• • •

Why It Matters

You can use Proton for:

- Real-time dashboards
- Streaming ETL and data prep
- Alerts and anomaly detection
- IoT and sensor data
- Log analytics and observability

It replaces a *stack* of tools with a single binary and a single language: SQL.

You can deploy it anywhere — Docker, Homebrew, cloud — or simply run it locally to experiment.

```
docker run -d -p 8123:8123 -p 8463:8463 \
--name proton d.timeplus.com/timeplus-io/proton:latest
```

• • •

My Take

Proton seems like a tool you try out of curiosity... and then realize it quietly replaces three things in your stack.

It's fast. It's clean. It doesn't yell at you for not setting a JVM flag, and it speaks the most universal language in data: SQL.

If you work with live data streams — or have ever wished Flink were easier — give this a try:

```
curl https://install.timeplus.com/oss | sh
```

You'll have a streaming SQL engine up and running before your coffee cools.

Sql

Sql Server

Live Streaming

AI

Data



Following ▾

Written by Algo Insights

3.4K followers · 6 following

Algo Insights: Unlocking the Future of Trading with Code, Data, and Intelligence.

https://linktr.ee/Algo_Insights

No responses yet



What are your thoughts?

More from Algo Insights

 In Coding Nexus by Algo Insights

4 Open-Source Tools That Made Me Rethink My Dev Setup

I've been coding for a while now. Most of the tools we use every day... they've been the same for years. Editors, browsers, frameworks. Just...

 Sep 3  450  9



 In Coding Nexus by Algo Insights

How to Build Your Own AI Rig for Running Local LLMs (Gemma, Mistral, Qwen, GPT-OSS and Llama)

About three months ago, I realised I was utterly dependent on companies that didn't care about anything except power, Money, and control.

 Oct 8  354  12



...

 In Coding Nexus by Algo Insights

DSPy: Stop Prompting, Start Programming Your AI

If you've ever spent hours tweaking prompts, you know how frustrating it can be.

 Oct 1  248  6



...

 Algo Insights

5 Prompting Secrets AI Engineers Use (That Most People Don't Know)

Most folks treat AI like Google. They type in a question, hit enter, and hope something useful comes back.

 Sep 25  376  10



...

See all from Algo Insights

Recommended from Medium

 In Meyka Finance Wire by Huzaifa Zahoor

How to Build a Stock Screener with Custom Technical Indicators

When a client needed to track 900 stocks with specific technical indicators, standard screening tools could not do the job. We built a...

 6d ago  87  3



...

 In Level Up Coding by Fareed Khan

Building an Agentic Deep-Thinking RAG Pipeline to Solve Complex Queries

Planning, Retrieval, Reflection, Critique, Synthesis and more

4d ago

801

4



...

In CodeToDeploy by TechToFit - Master Your Life with Tech

I Tried Google's New AI Agents. It's a Gold Rush.

I spend my days deep in the world of AI, but every so often, something drops that makes me stop everything. This is one of those times...

Oct 10

164

3



...

Algo Insights

10 AI Websites You Can't Ignore in 2025

AI has sped up this year, reminding me of times I'd spend hours pounding on a keyboard to fix a bug.

◆ 6d ago ⌘ 13 🗣 1



...

◻ In Towards Deep Learning by Sumit Pandey

Why Everyone Will Want DGX Spark on Their Desk—Yes, Everyone

I just saw this picture today and was amazed, I've been waiting for this moment for a long time. (No, it's not Elon.) It's that tiny...

◆ Oct 14 ⌘ 23 🗣 7



...

 In Coding Nexus by Civil Learning

MarkItDown: Convert Anything into Markdown—the Smart Way to Feed LLMs

You know that feeling when you're trying to feed a PDF or a Word document into an LLM, and it just doesn't understand what's going on...

 Oct 15  144  3



...

[See more recommendations](#)