

T3CH · [Follow publication](#)

Cloud-Native RAG: embeddings + vectors, fully self-hosted

15 min read · 1 day ago



using System;

[Follow](#)

Listen

Share

More



An agent without a pre-trained model needs your data. You inject your documents into a vector database the agent can query to answer questions or drive actions for your specific needs.

We will see in this new article how to host your RAG solution (embeddings + vector DB) fully self-hosted in an industrial way. Then we'll actually use it to implement RAGs and connect your agents. Examples are in .NET. Sorry to the Python lovers — Python doesn't have a monopoly on AI 😊

The principle is basic: to create an embedding, one POST is enough, and integration with QDrant is straightforward via SDKs across stacks (Python, .NET, Java, Rust...). For agents (LangChain, Semantic Kernel...), all integrate natively with QDrant.


As usual, deployment scripts are done with Terraform to deploy Kubernetes resources.


We'll use:

- [KServe](#) to host our embedding inference model.
- [Envoy AI Gateway](#) to securely and uniformly expose the inference model.
- [QDrant](#) for the vector database.
- [.NET 10](#) to implement a rag tool to index documents and mcp tool for the agent.

All of it deployed on a [Kubernetes](#) cluster.

High-Level architecture diagram





Why self-host?

Starting in **SaaS** for embeddings (OpenAI, Anthropic, Mistral...) is the right way. GPU instances (even though embeddings can run on CPU) have a steep fixed cost and there's infrastructure to maintain.

But when **costs explode** or when you have **sovereignty constraints** for your data, **self-hosted** becomes the way to go.

Kubernetes and the **open-source** community make cloud-native AI more accessible than ever. That's exactly what we'll build together here.

Step 1 — Install KServe CRDs

Prereqs: Kubernetes > 1.30, cert-manager (> 1.15.0) on your cluster.

Why install CRDs separately?

Always preferable to install CRDs in a **common infrastructure layer**, and install implementations in separate terraform layers.

```

resource "helm_release" "kserve" {

  name          = "kserve-crd"
  namespace     = "kube-system"
  create_namespace = false
  chart         = "kserve-crd"
  repository    = "oci://ghcr.io/kserve/charts"
  version       = "v0.16.0" # Verion of 2025 october
}

```

CRDs deployed

```
using system@% kubectl get crd | grep kserve
clusterservingruntimes.serving.kserve.io
clusterstoragecontainers.serving.kserve.io
inferencegraphs.serving.kserve.io
inferenceservices.serving.kserve.io
localmodelcaches.serving.kserve.io
localmodelnodegroups.serving.kserve.io
localmodelnodes.serving.kserve.io
servingruntimes.serving.kserve.io
trainedmodels.serving.kserve.io

2025-07-24T18:21:5
2025-07-24T18:21:5
2025-07-24T18:21:5
2025-07-24T18:21:5
2025-07-24T18:21:5
2025-07-24T18:21:5
2025-07-24T18:21:5
2025-07-24T18:21:5
2025-07-24T18:21:5
```

What they do :

- **InferenceService** – classic KServe resource for deploying a predictive/generative model behind a stable URL. Defines predictor (runtime, model, resources), ingress, and autoscaling.
- **ClusterStorageContainer** – cluster-level definition of model sources (OCI, model kits, HF, S3) and the init-container that downloads/extracts models. Other CRDs reference it via storage config.
- **LocalModelNodeGroup** – labels a set of nodes where KServe may pre-cache models locally; sets disk quotas for cache.
- **LocalModelNode** – per-node state of the local cache (downloaded/failed/not ready). Controlled by KServe for fine-grained troubleshooting.
- **LocalModelCache** – declares that a given model must be pre-downloaded to specific node groups (reducing cold-start latency).

Step 2 — Install KServe (Standard mode)

```
resource "kubernetes_namespace" "kserve" {
  metadata {
    name = "kserve"

    labels = {
      provisioned_by = "terraform"
    }
  }
}
```

```

        }

    }

resource "helm_release" "kserve" {

  name      = "kserve"
  namespace = kubernetes_namespace.kserve.metadata[0].name
  chart     = "kserve"
  repository = "oci://ghcr.io/kserve/charts"
  version   = "v0.16.0" # Version of 2025 october

  values = [
    yamlencode({
      kserve = {
        controller = {
          = "Standard"
        }
        storage = {
          resources = {
            requests = {
              cpu      = "100m"
              memory = "2Gi"
            }
            limits = {
              cpu      = "1"
              memory = "8Gi"
            }
          }
        }
      }
    })
  ]
}

```

Standard vs Knative

Standard → KServe creates vanilla Kubernetes resources (Deployment, Service, HPA/KEDA, Ingress). Recommended for prod and LLMs if you don't need scale-to-zero.

Knative → Serverless (scale-to-zero, revisions, canary). Powerful, but requires Knative (+ data plane like Istio) and currently **doesn't work natively with Envoy AI Gateway** in our path.

Storage resources matter

storage sets resources for the **init-container** that downloads models. Defaults are often too low for multi-GB models → pod crashes. You can override per backend (HF/S3/etc.).

All chart values here.

Validate KServe install :

```
using system@~ % kubectl get all -n kserve
NAME                                     READY   STATUS    R
pod/kserve-controller-manager-65b6cc8469-szx2w   2/2     Running   0
pod/kserve-localmodel-controller-manager-5b4778d898-l5562   1/1     Running   0

NAME                           TYPE      CLUSTER-IP        EXTER
service/kserve-controller-manager-service   ClusterIP   XXX.XXX.XX.XX    <none>
service/kserve-webhook-server-service       ClusterIP   XXX.XXX.XXX.XXX   <none>
```

Classic CRD/controller pattern: KServe's controller watches your KServe manifests and creates the needed Kubernetes resources.

Step 3 — Choose the embedding model

We'll use **bge-m3**. It's a solid CPU-mode starting point: multilingual, versatile, easy to deploy.

If your top priority is absolute top scores and you can allocate a GPU, **gte-Qwen2-7B-instruct** or **jina-v3** can do better.

We'll proceed with **bge-m3**, downloaded from [Hugging Face](#):

≈ 3 GB. If you only need English, consider **bge-large-en-v1.5**, **bge-base-en-v1.5**, or smaller **bge-small-en-v1.5**.

Family table :

Model Name	Dimension	Sequence Length	Introduction
BAAI/bge-m3	1024	8192	multilingual; unified
BAAI/bge-large-en-v1.5	1024	512	English model
BAAI/bge-base-en-v1.5	768	512	English model
BAAI/bge-small-en-v1.5	384	512	English model

Step 4 — Host bge-m3 on KServe and expose REST

```

resource "kubectl_manifest" "inference_service" {

    yaml_body = yamlencode({
        apiVersion = "serving.kserve.io/v1beta1"
        kind       = "InferenceService"
        metadata   = {
            name      = "bge-m3"
            namespace = "using-system-models"
            annotations = {
                "serving.kserve.io/autoscalerClass" = "hpa"
            }
        }
        spec = {
            predictor = {
                minReplicas = 2
                maxReplicas = 3
                scaleTarget = 80
                scaleMetric = "cpu"

                model = {
                    modelFormat = {
                        name = "huggingface"
                    }
                    args = [
                        "--model_name=bge-m3",
                        "--backend=huggingface",
                        "--task=text_embedding",
                    ]
                    storageUri = "hf://BAAI/bge-m3"

                    resources = {
                        requests = {
                            cpu      = "2"
                            memory  = "12Gi"
                        }
                        limits = {
                            cpu      = "4"
                            memory  = "24Gi"
                        }
                    }
                }
            }
        }
    })
}

```

We host fully on CPU (no GPU). Add `tolerations` and `nodeSelector` under `predictor` to land on the right nodes.

Autoscaling :

```
"serving.kserve.io/autoscalerClass" = "hpa"

minReplicas = 2
maxReplicas = 3
scaleTarget = 80
scaleMetric = "cpu"
```

This scales 2→3 replicas with an 80% CPU trigger. HPA is simple; for metric-based scaling (e.g., workflow requests) use **KEDA**. For scale-to-zero, use **Knative**.

Model args :

```
modelFormat = { name = "huggingface" }
args = [
    "--model_name=bge-m3",
    "--backend=huggingface",
    "--task=text_embedding",
]
storageUri = "hf://BAAI/bge-m3"
```

- Use HF backend, set task `text_embedding`, download from `hf://BAAI/bge-m3`.

Kubernetes resources created :

```
using system@~ % kubectl get all -n using-system-models
NAME                                     READY   STATUS    RESTARTS   AGE
pod/bge-m3-predictor-6d6f6c49c7-7skmt   2/2     Running   0          84m
pod/bge-m3-predictor-6d6f6c49c7-z2v9z   2/2     Running   0          84m

NAME           TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)
service/bge-m3-predictor   ClusterIP   XXX.XXX.XXX.XX  <none>        80/TCP

NAME           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/bge-m3-predictor  2/2     2           2          3d18h
```

NAME	REFERENCE
horizontalpodautoscaler.autoscaling/bge-m3-predictor	Deployment/bge-m3-predic

Test with curl :

```
curl -X POST "http://bge-m3-predictor.using-system-models.svc.cluster.local/ope
-H "Content-Type: application/json" \
-H "Accept: application/json" \
-H "Accept-Encoding: identity" \
-d '{
  "model": "bge-m3",
  "input": "This is just a basic embedding test."
}'
```

Response :

```
{
  "id": "0ebbb54d-802b-4af1-8738-f28f770e194f",
  "object": "embedding",
  "created": 1761117146,
  "model": "bge-m3",
  "data": [
    {
      "index": 0,
      "object": "embedding",
      "embedding": [
        -0.01863192766904831,
        0.0039021808188408613,
        -0.011211015284061432,
        -0.00893788319081068,
        0.009318851865828037,
        ...
        0.020269472151994705
      ]
    }
  ],
  "usage": {
    "prompt_tokens": 12,
    "total_tokens": 12,
    "completion_tokens": 0,
    "prompt_tokens_details": null
  }
}
```

```
    }
}
```

Pod init can be slow (model download). We'll cover **KServe Local Model Cache** in bonus to slash cold-start times.

Step 5 — Add Envoy AI Gateway (security, routing, observability)

Not mandatory for simple needs, but crucial for auth, rate limiting, observability, fallbacks, federating models.

- Doc: <https://aigateway.envoyproxy.io/>
- Install guide : <https://usingsystem.io/install-an-ai-gateway-with-envoy-gateway-ai-7988ac28f901>
- OpenAI backend example : <https://usingsystem.io/openai-backend-with-envoy-ai-gateway-3cc4c438effb>

We'll expose KServe bge-m3 almost like a SaaS provider.

We need **AIServiceBackend**, **ReferenceGrant**, and **AIGatewayRoute**.

```
variable "name" {
  description = "The name of the model to be deployed."
  type        = string
  default     = "bge-m3"
}

variable "namespace" {
  description = "The namespace in which to deploy the model."
  type        = string
}

variable "ai_gateway_namespace" {
  description = "The namespace for the AI Gateway."
  type        = string
}

variable "ai_gateway_name" {
  description = "The name of the AI Gateway."
  type        = string
}
```

```

locals {
    ai_backend_name      = "${var.name}-predictor"
    ai_backend_service_name = "ai-backend-service-${var.name}"
}

resource "kubectl_manifest" "ai_service_backend" {

    depends_on = [
        kubectl_manifest.inference_service
    ]

    yaml_body = yamlencode({
        apiVersion = "aigateway.envoyproxy.io/v1alpha1"
        kind       = "AIServiceBackend"
        metadata   = {
            name      = local.ai_backend_service_name
            namespace = var.ai_gateway_namespace
        }
        spec      = {
            schema = {
                name = "OpenAI"
            }
            backendRef = {
                name      = local.ai_backend_name
                namespace = var.namespace
                kind      = "Service"
                group    = ""
                port     = 80
            }
            timeouts = {
                request = "300s"
            }
        }
    })
}

resource "kubectl_manifest" "ai_service_backend_grant" {

    depends_on = [
        kubectl_manifest.ai_service_backend
    ]

    yaml_body = yamlencode({
        apiVersion = "gateway.networking.k8s.io/v1beta1"
        kind       = "ReferenceGrant"
        metadata   = {
            name      = "envoy-gateway-ref-grant-${var.name}"
            namespace = var.namespace
        }
        spec      = {
            from = [
                {
                    group      = "gateway.networking.k8s.io"

```

```

        kind      = "HTTPRoute"
        namespace = var.ai_gateway_namespace
    }
]
to = [
{
    group = ""
    kind   = "Service"
}
]
}
)
}

resource "kubectl_manifest" "ai_backend_route" {

depends_on = [
kubectl_manifest.inference_service,
kubectl_manifest.ai_service_backend,
kubectl_manifest.ai_service_backend_grant
]

yaml_body = yamlencode({
apiVersion = "aigateway.envoyproxy.io/v1alpha1"
kind       = "AIGatewayRoute"
metadata = {
name      = "envoy-ai-gateway-${var.name}-route"
namespace = var.ai_gateway_namespace
}
spec = {
schema = {
name = "OpenAI"
}
parentRefs = [
{
name      = var.ai_gateway_name
namespace = var.ai_gateway_namespace
kind      = "Gateway"
group     = "gateway.networking.k8s.io"
}
]
rules = [
{
matches = [
{
headers = [
{
type  = "Exact"
name   = "x-ai-eg-model"
value  = var.name
}
]
}
]
}
]
}
}
}
}
```


```

        ]
        backendRefs = [
            {
                name    = local.ai_backend_service_name
                weight = 100
            }
        ]
    }
]
llmRequestCosts = [
{
    metadataKey = "llm_input_token"
    type       = "InputToken"
},
{
    metadataKey = "llm_output_token"
    type       = "OutputToken"
},
{
    metadataKey = "llm_total_token"
    type       = "TotalToken"
}
]
}
)
}
]

```

What this does

- **AIServiceBackend** — declares the backend with **OpenAI** schema and points to `Service/bge-m3-predictor:80`.
- **ReferenceGrant** — authorizes a cross-namespace **HTTPRoute** to reference that **Service**.
- **AIGatewayRoute** — attaches to your **Gateway**, matches header `x-ai-eg-model: bge-m3`, routes to the backend, and records token usage via `llmRequestCosts`.



```
v  
[ Service/bge-m3-predictor:80 ]
```

KServe URL prefix tweak (important for Envoy AI Gateway integration) :

```
resource "kubectl_manifest" "inference_service" {  
  
    yaml_body = yamlencode({  
        apiVersion = "serving.kserve.io/v1beta1"  
        kind       = "InferenceService"  
        metadata   = {  
            name      = "bge-m3"  
            namespace = "using-system-models"  
            annotations = {  
                "serving.kserve.io/autoscalerClass" = "hpa"  
            }  
        }  
        spec     = {  
            predictor = {  
                ....  
                env  = [  
                    {  
                        name  = "KSERVE_OPENAI_ROUTE_PREFIX",  
                        value = ""  
                    }  
                ]  
            }  
        }  
    })  
}
```

Discover models via Envoy :

```
curl -X GET "http://envoy-ai-gateway.envoy-gateway-system.svc.cluster.local/mod  
-H "Content-Type: application/json" \  
-H "Accept: application/json"
```

```
{  
  "object": "list",
```

```

"data": [
  {
    "id": "gpt-4",
    "object": "model",
    "created": 1687882411,
    "owned_by": "Envoy AI Gateway"
  },
  {
    "id": "bge-m3",
    "object": "model",
    "created": 1687882411,
    "owned_by": "Envoy AI Gateway"
  },
  ...
]
}

```

Query embeddings via Envoy :

```

curl -X POST "http://envoy-ai-gateway.envoy-gateway-system.svc.cluster.local/v1
-H "Content-Type: application/json" \
-H "Accept: application/json" \
-H "Accept-Encoding: identity" \
-d '{
  "model": "bge-m3",
  "input": "This is just a basic embedding test."
}'

```

Step 6 — Install QDrant (vector DB)

Qdrant is a Rust vector DB for high-dimensional similarity. In cluster mode, it shards and replicates for HA. The Helm chart enables distributed mode by bumping `replicaCount` and setting `config.cluster.enabled: true`. Production hardening: storage, network, anti-affinity, observability, upgrade strategies.

Terraform :

```

resource "kubernetes_namespace" "qdrant" {
  metadata {
    name = "qdrant"

    labels = {

```

```
        provisioned_by  = "terraform"
    }
}
}

resource "helm_release" "qdrant" {

  name      = "qdrant"
  namespace = kubernetes_namespace.qdrant.metadata[0].name
  chart     = "qdrant"
  repository = "https://qdrant.github.io/qdrant-helm"
  version   = "1.15.5" // Version en date d'octobre 2025

  values = [
    yamlencode({
      replicaCount = 3

      resources = {
        requests = {
          cpu      = "2"
          memory = "12Gi"
        }
        limits = {
          cpu      = "4"
          memory = "24Gi"
        }
      }
    }

    persistence = {
      enabled = true
      size    = "30Gi"
    }

    podDisruptionBudget = {
      enabled        = true
      maxUnavailable = 1
    }

    affinity = {
      podAntiAffinity = {
        preferredDuringSchedulingIgnoredDuringExecution = [
          {
            weight = 100
            podAffinityTerm = {
              topologyKey = "kubernetes.io/hostname"
              labelSelector = {
                matchLabels = { "app.kubernetes.io/name" = "qdrant" }
              }
            }
          }
        ]
      }
    }
  ]
}
```

```

        config = {
            cluster = {
                enabled = true
            }
        }
    })
]
}

```

- **PodDisruptionBudget (PDB)** — `maxUnavailable = 1`:
Limits voluntary disruptions (node drains, upgrades, cluster autoscaler evictions) so that at most one Qdrant pod can be down at a time. With `replicaCount >= 2`, Kubernetes will serialize evictions to keep the service available. This is particularly helpful during rolling upgrades of nodes or when draining for maintenance.
- **Pod anti-affinity (preferred) across hosts:**
`preferredDuringSchedulingIgnoredDuringExecution` with `topologyKey = "kubernetes.io/hostname"` nudges the scheduler to spread Qdrant pods across different nodes. This reduces correlated failure risk (one node loss shouldn't take out all replicas).
Using **preferred** (instead of **required**) keeps scheduling flexible when capacity is tight. If you need strict separation, you can switch to `requiredDuringSchedulingIgnoredDuringExecution`—but be aware it may block scheduling in small clusters.

Services :

```

using system@~ % kubectl get svc -n qdrant -o wide
NAME           TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)
qdrant         ClusterIP  XXX.XXX.XX.XXX  <none>          6333/TCP,6334/TCP,
qdrant-headless   ClusterIP  None           <none>          6333/TCP,6334/TCP,

```

CPU vs GPU

For **query**, GPU won't change much. Gains are mainly for **heavy indexing** with large embeddings (bge-m3 is 1024-D), where GPU can bring notable speedups (often ~10x).

Step 7 — Development: .NET end-to-end

To index your documents, create an embedding and send vectors into QDrant.

Create embedding with a single POST :

```
var embeddingApiUrl = "http://envoy-ai-gateway.envoy-gateway-system.svc.cluster.local";
var httpClient = new HttpClient();

async Task<List<float>> CreateEmbeddingAsync(string text)
{
    // Build JSON manually to avoid warnings
    var escapedText = text.Replace("\\\\", "\\\\\\\").Replace("\\\"", "\\\"");
    var requestBody = $"{{{{\"model\": \"bge-m3\", \"input\": \"{escapedText}\",}}}}";

    var content = new StringContent(requestBody, System.Text.Encoding.UTF8, "application/json");
    var response = await httpClient.PostAsync(embeddingApiUrl, content);
    response.EnsureSuccessStatusCode();

    var responseBody = await response.Content.ReadAsStringAsync();
    using var embeddingDoc = JsonDocument.Parse(responseBody);

    // Extract embedding vector from response
    var embeddingArray = embeddingDoc.RootElement
        .GetProperty("data")[0]
        .GetProperty("embedding");

    var vectors = new List<float>();
    foreach (var value in embeddingArray.EnumerateArray())
    {
        vectors.Add((float)value.GetDouble());
    }

    return vectors;
}
```

Qdrant client (Nuget [Qdrant.Client](#)) :

```
var qdrantHost = "qdrant.qdrant.svc.cluster.local";
var qdrantClient = new QdrantClient(qdrantHost);
```

Create (or recreate) a collection :

```

async Task CreateOrRecreateCollectionAsync(string collectionName, uint vectorSize)
{
    // Check if collection exists and delete if it does
    var collections = await qdrantClient.ListCollectionsAsync();
    if (collections.Contains(collectionName))
    {
        await qdrantClient.DeleteCollectionAsync(collectionName);
    }

    // Create collection
    await qdrantClient.CreateCollectionAsync(
        collectionName: collectionName,
        vectorsConfig: new VectorParams { Size = vectorSize, Distance = DistanceType.Cosine });
}

```

Insert an embedding :

```

await CreateOrRecreateCollectionAsync("myCollection", 1024);
...
string text = "...";
string language = "en-US";

var embedding = await CreateEmbeddingAsync(text);

    var point = new PointStruct
    {
        Id = Guid.NewGuid(),
        Vectors = embedding.ToArray(),
        Payload =
        {
            ["text"] = text,
            ["language"] = language ?? "en-us"
            ...
        }
    };

await qdrantClient.UpsertAsync("myCollection", new[] { point });

```

When your agent queries the vector DB, it retrieves nearest vectors and the payload becomes the context. Structure the payload thoughtfully.

Agent integration

Two options :

- Use native Qdrant integrations in your agent framework if available.
- Or keep it simple via a **Tool** or a **Model Context Protocol (MCP)** server.

In this article we will choose the option 2.

Service building context from Qdrant :

```
using System.Text;
using System.Text.Json;
using Qdrant.Client;

namespace ms_doc_assistant.Services;

public class QdrantRagService
{
    private readonly QdrantClient _qdrantClient;
    private readonly HttpClient _embeddingHttpClient;
    private readonly string _embeddingApiUrl;
    private readonly string _embeddingModel;

    public QdrantRagService(
        string qdrantHost,
        string embeddingApiUrl,
        string embeddingModel)
    {
        _qdrantClient = new QdrantClient(qdrantHost);
        _embeddingHttpClient = new HttpClient();
        _embeddingApiUrl = embeddingApiUrl;
        _embeddingModel = embeddingModel;
    }

    /// <summary>
    /// Build context string from search results for the agent
    /// </summary>
    public async Task<string> BuildContextForQueryAsync(string query, int limit)
    {
        var documents = await SearchDocumentsAsync(query, limit);

        if (documents.Count == 0)
        {
            return "No relevant documents found.";
        }

        var contextBuilder = new StringBuilder();
        contextBuilder.AppendLine("### Relevant Knowledge Base Articles:");
        contextBuilder.AppendLine();
    }
}
```

```

        foreach (var (title,
                      category,
                      tags,
                      summary,
                      content,
                      score) in documents)
    {
        contextBuilder.AppendLine($"**{title}** ({category}) [relevance: {score}]");

        if (tags.Length > 0)
        {
            contextBuilder.AppendLine($"Tags: {string.Join(", ", tags)}");
        }

        if (!string.IsNullOrWhiteSpace(summary))
        {
            contextBuilder.AppendLine($"Summary: {summary}");
        }

        if (content.Length > 0)
        {
            contextBuilder.AppendLine("Key Points:");
            foreach (var section in content)
            {
                contextBuilder.AppendLine($" - {section}");
            }
        }

        contextBuilder.AppendLine();
    }

    return contextBuilder.ToString();
}

private async Task<List<(string title,
                           string category,
                           string[] tags,
                           string summary,
                           string[] content,
                           double score)>>
SearchDocumentsAsync(
    string query,
    int limit = 5)
{
    var queryEmbedding = await CreateEmbeddingAsync(query);

    var results = await _qdrantClient.SearchAsync(
        collectionName: "knowledge_base",
        vector: queryEmbedding.ToArray(),
        limit: (ulong)limit
    );
}

```

```

        return results.Select(r => (
            title: r.Payload["title"].StringValue,
            category: r.Payload["category"].StringValue,
            tags: r.Payload["tags"].ListValue.Values.Select(v => v.StringValue)
            summary: r.Payload["summary"].StringValue,
            content: r.Payload["content"].ListValue.Values.Select(v => v.StringValue)
            score: (double)r.Score
        )).ToList();
    }

    /// <summary>
    /// Create embedding for text using the embedding API
    /// </summary>
    private async Task<List<float>> CreateEmbeddingAsync(string text)
    {
        var escapedText = text.Replace("\\\\", "\\\\").Replace("\\\"", "\\\"");
        var requestBody = $"\"{{\"model\":\"{_embeddingModel}\",\"input\":\"{esca

        var content = new StringContent(requestBody, Encoding.UTF8, "application/json");
        var response = await _embeddingHttpClient.PostAsync(_embeddingApiUrl, content);
        response.EnsureSuccessStatusCode();

        var responseBody = await response.Content.ReadAsStringAsync();
        using var embeddingDoc = JsonDocument.Parse(responseBody);

        var embeddingArray = embeddingDoc.RootElement
            .GetProperty("data")[0]
            .GetProperty("embedding");

        var vectors = new List<float>();
        foreach (var value in embeddingArray.EnumerateArray())
        {
            vectors.Add((float)value.GetDouble());
        }

        return vectors;
    }
}

```

Expose as an MCP Tool (via nuget [ModelStateProtocol.AspNetCore](#)) :

```

using System.ComponentModel;
using ModelStateProtocol.Server;
using ms_doc_assistant.Services;

namespace ms_doc_assistant.Tools;

[McpServerToolType]

```

```

public class RagTool(QdrantRagService ragService)
{
    [McpServerTool(Name = "SearchKnowledgeBase")]
    [Description("Search for relevant technical documents in the knowledge base")]
    public async Task<string> SearchKnowledgeBase(
        [Description("The search query for the internal documentation or proced
    {
        return await ragService.BuildContextForQueryAsync(query);
    }
}

```

And now the bonus tracks !

Bonus #1 — Complete RAG loader

```

#:package Qdrant.Client@1.12.0

using System.Text.Json;
using Qdrant.Client;
using Qdrant.Client.Grpc;

var embeddingApiUrl = Environment.GetEnvironmentVariable("EMBEDDING_API_URL") ??
var embeddingModel = Environment.GetEnvironmentVariable("EMBEDDING_MODEL") ?? ""
var qdrantHost = Environment.GetEnvironmentVariable("QDRANT_HOST") ?? "qdrant.q
var datasetsPath = Environment.GetEnvironmentVariable("DATASETS_PATH") ?? "../.
var datasetsSearchPattern = Environment.GetEnvironmentVariable("DATASETS_SEARCH

// Initialize HTTP client for embeddings
var httpClient = new HttpClient();

// Initialize Qdrant client
var qdrantClient = new QdrantClient(qdrantHost);

// Create or recreate collection
await CreateOrRecreateCollectionAsync("knowledge_base", 1024);

var jsonFiles = Directory.GetFiles(datasetsPath, datasetsSearchPattern);
Console.WriteLine($"Number of files found: {jsonFiles.Length}\n");

foreach (var filePath in jsonFiles)
{
    Console.WriteLine($"Processing file: {Path.GetFileName(filePath)}");
    using var document = JsonDocument.Parse(File.ReadAllText(filePath));

    var documents = document.RootElement.EnumerateArray().ToList();

    await Parallel.ForEachAsync(documents, new ParallelOptions { MaxDegreeOfPar
    {

```

```

var title = doc.GetProperty("title").GetString();
var category = doc.GetProperty("category").GetString();
var tags = doc.GetProperty("tags");
var summary = doc.GetProperty("summary");
var content = doc.GetProperty("content");

Console.WriteLine($"\\n Processing Document: {title}");

if (string.IsNullOrWhiteSpace(title))
    return;

var titleEmbedding = await CreateEmbeddingAsync(title);

var point = new PointStruct
{
    Id = Guid.NewGuid(),
    Vectors = titleEmbedding.ToArray(),
    Payload =
    {
        ["title"] = title,
        ["category"] = category ?? "general",
        ["tags"] = tags.EnumerateArray().Select(t => t.GetString()!).To
        ["summary"] = summary.GetString() ?? "",
        ["content"] = content.EnumerateArray().Select(s => s.GetString()
    }
};

await qdrantClient.UpsertAsync("knowledge_base", new[] { point });
}
}

Console.WriteLine("\\nProcessing completed!");

```

#region Helper Methods

```

// Create embedding from text using embedding API
async Task<List<float>> CreateEmbeddingAsync(string text)
{
    var escapedText = text.Replace("\\", "\\\\").Replace("'", "\\'");
    var requestBody = $"\"{\"model\": \"{embeddingModel}\", \"input\": \"{escapedTe

    var content = new StringContent(requestBody, System.Text.Encoding.UTF8, "ap
    var response = await httpClient.PostAsync(embeddingApiUrl, content);
    response.EnsureSuccessStatusCode();

    var responseBody = await response.Content.ReadAsStringAsync();
    using var embeddingDoc = JsonDocument.Parse(responseBody);

    var embeddingArray = embeddingDoc.RootElement
        .GetProperty("data")[0]
        .GetProperty("embedding");

    var vectors = new List<float>();

```

```

        foreach (var value in embeddingArray.EnumerateArray())
        {
            vectors.Add((float)value.GetDouble());
        }

        return vectors;
    }

    // Create or recreate a Qdrant collection
    async Task CreateOrRecreateCollectionAsync(string collectionName, uint vectorSize)
    {
        var collections = await qdrantClient.ListCollectionsAsync();
        if (collections.Contains(collectionName))
        {
            Console.WriteLine($"Collection '{collectionName}' already exists, deleting it");
            await qdrantClient.DeleteCollectionAsync(collectionName);
            Console.WriteLine($"Collection '{collectionName}' deleted\n");
        }

        await qdrantClient.CreateCollectionAsync(
            collectionName: collectionName,
            vectorsConfig: new VectorParams { Size = vectorSize, Distance = DistanceType.Cosine });
        Console.WriteLine($"Collection '{collectionName}' created successfully\n");
    }
}

#endregion

```

Bonus #2 — KServe Local Model Cache (reduce cold-starts drastically)

Key idea: enable `localmodel` in KServe → a DaemonSet coordinates per-node caches.
You define:

- **ClusterStorageContainer** (how to download `hf://` models),
- **LocalModelNodeGroup** (where to store on nodes; PVC/hostPath + node affinity),
- **LocalModelCache** (which model to pre-download and to which node groups).

Your manifests (HF secret, storage initializer, NodeGroup, LocalModelCache, reconciler DaemonSet) are kept exactly as provided in your original text, including the tips about new nodes / SPOT nodes and the annotation trick to force reconciliation.

After that: delete your predictor pods → restart is **much faster** (model already on disk). If downloads fail, check logs in `kserve-localmodel-jobs`.

Conclusion

Self-hosting a cloud-native RAG stack is not only possible — it's clean:

- KServe serves **bge-m3** for embeddings (CPU-friendly),
- Envoy AI Gateway unifies access with an OpenAI-style API plus auth, quotas, and metrics,
- Qdrant stores vectors with strong recall/latency and simple operations.

Links

- [KServe](#)
- [KServe chart values](#)
- [Envoy AI Gateway](#)
- [Qdrant](#)
- [Qdrant Helm](#)
- [Hugging Face bge-m3](#)
- [Install Envoy AI Gateway](#)
- [OpenAI backend with Envoy](#)
- [NuGet Qdrant.Client](#)
- [NuGet Aspire.Qdrant.Clien](#)
- [NuGet ModelContextProtocol.AspNetCore](#)
- [Knative](#)
- [KEDA](#)
- [cert-manager](#)

Benchmarked **bge-m3** vs **gte-Qwen2-7B-instruct** on your data? Tuned Qdrant HNSW for 1024-D? Hit quirks with Envoy ↔ KServe pathing? Share your configs, metrics, and war stories.

Kubernetes

Mlops

Envoy Proxy

Terraform

AI



Follow

Published in T3CH

1.3K followers · Last published 7 hours ago

Snoop & Learn about Technology, AI, Hacking, Coding, Software, News, Tools, Leaks, Bug Bounty, OSINT & Cybersecurity !! But, not limited 2, anything that is Tech Linked...You'll probably find here !;)—Stay ahead with Latest Tech News! -> You write about? Just ping to join !



Follow



Written by using System;

157 followers · 103 following

Medium blog from Matthieu Vlad, .NET Senior Dev, DevOps & Cloud Architect. Support me by send your tips : usingsystem.eth ❤️

No responses yet



What are your thoughts?

More from using System; and T3CH

In T3CH by using System;

Bitnami's Keycloak Shake-Up: What Broke, How to Buy Time, and Where to Migrate (with Terraform +...)

This post is a follow-up to my earlier article, “Open-source and cloud-native identity management and gateway”.

Sep 18  97  1



...

In T3CH by Frost

How I Instantly Know You Used ChatGPT

These days, AI tools are everywhere, and ChatGPT is leading the pack. It's easy, and it's tempting to use for just about everything: school...

★ Sep 9 ⚡ 169 🗣 6



...

In T3CH by Ali Hamza

Mastering Docker on WSL2: A Complete Guide Without Docker Desktop

Streamline Your Docker Experience on Windows with WSL2—No Docker Desktop Needed!

Jan 7 ⚡ 188 🗣 7



...

In T3CH by using System;

Deploying Your AKS Cluster with Terraform: Key Points for a Successful Production Rollout

Azure Kubernetes Service (AKS) is Microsoft Azure's fully managed Kubernetes offering. It simplifies many operational complexities around...

Dec 29, 2024  62  1



...

See all from using System;

See all from T3CH

Recommended from Medium

 In Data, AI and Beyond by Julius Nyerere Nyambok

Generate AWS architectural diagrams using this simple method

A combination of LLMs and MCP Servers.

 Oct 14  64  1



...

 DhanushKumar

LLM poisoning

What is LLM poisoning ?

5d ago  278  12



...

 In Netflix TechBlog by Netflix Technology Blog

Behind the Streams: Real-Time Recommendations for Live Events Part 3

By: Kris Range, Ankush Gulati, Jim Isaacs, Jennifer Shin, Jeremy Kelly, Jason Tu

6d ago  91  1



 In AWS Tip by Osman ALP

Mastering Kubernetes Observability: Deploying Grafana with Loki, Tempo, Mimir, and Alloy

The Modern Observability Trinity: Metrics, Logs, Traces, and Telemetry Collection

 Oct 13  10



...

 In AWS in Plain English by Aman Pathak | DevOps | AWS | K8s | Terraform

AWS Outage Root Cause Revealed: How a DNS Mismatch Affected DynamoDB and Core Services

Discover what caused the October 2025 AWS us-east-1 outage—a DNS mismatch that broke DynamoDB and triggered cascading failures across...

 6d ago  68  4



...

 In Opensearch Stories by Franco martin

Running OpenSearch in AWS

Is it any good?

Oct 7  1



...

[See more recommendations](#)