Member-only story

The 10 AI Agent Platforms Revolutionizing Software in 2025

12 min read · Oct 8, 2025

Micheal Lanham Following

Listen Share More

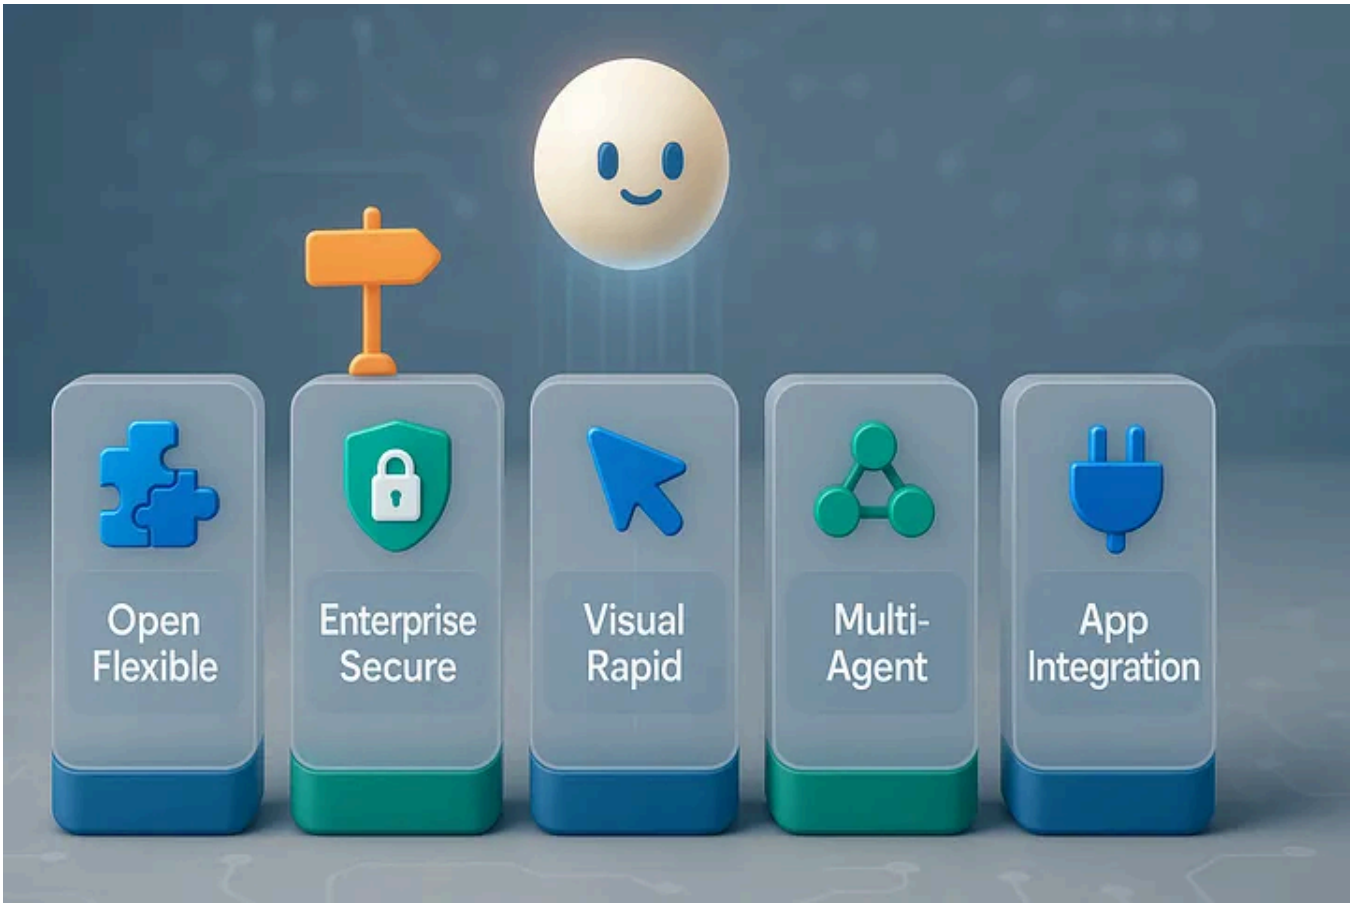

all images generated with gpt-image-1

A hands-on guide to choosing the right framework for building autonomous AI agents — from open-source powerhouses to enterprise solutions

You’re here because you’ve seen the future, and it’s autonomous. AI agents aren’t just chatbots anymore — they’re sophisticated programs that analyze information, make decisions, and execute complex tasks without constant human babysitting. And honestly? The explosion in capabilities by 2025 has been nothing short of mind-blowing.

I’ve spent the last few months diving deep into the AI agent ecosystem, and I want to share what I’ve learned. Whether you’re a startup founder looking to build your first AI assistant or an enterprise architect evaluating platforms, this guide breaks down the top 10 general-purpose AI agent platforms that are shaping how we build intelligent software.

TL;DR: The AI agent landscape has matured dramatically. LangChain dominates with flexibility, OpenAI and Google are pushing enterprise-grade SDKs, Microsoft offers collaboration-focused frameworks, while platforms like Dify democratize agent development with no-code tools. Your choice depends on three factors: technical complexity you’re willing to handle, ecosystem lock-in you can accept, and scale of deployment you’re targeting.

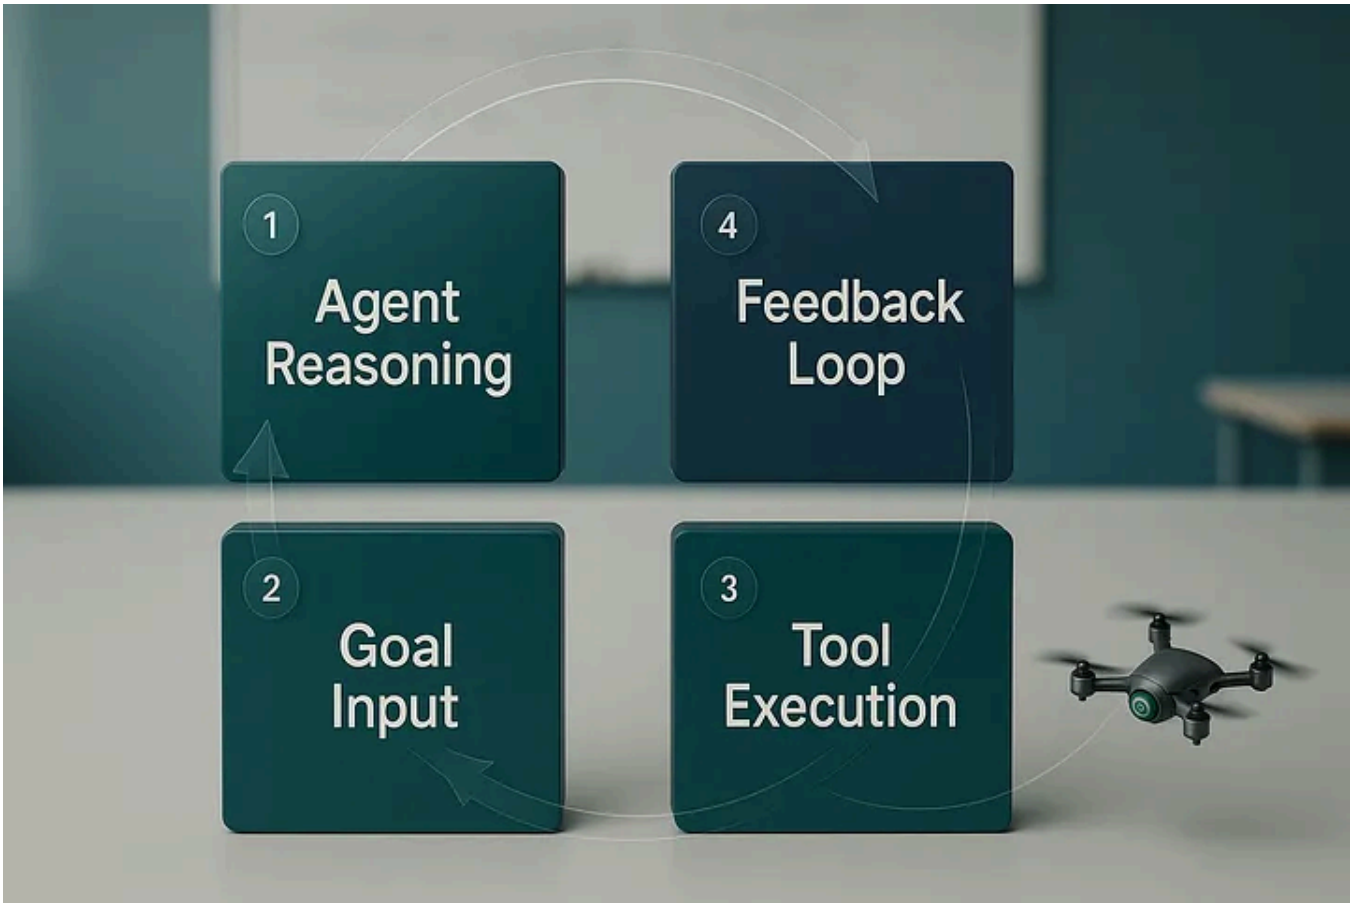

Understanding AI Agents: The Foundation

Before we dive into platforms, let's get clear on what we mean by "AI agents."

Think of an AI agent as an autonomous problem-solver. Unlike traditional software that follows predetermined paths, agents can:

- Break down complex goals into sub-tasks
- Choose and use appropriate tools
- Learn from feedback and iterate
- Maintain context across long interactions

The magic happens when you combine powerful LLMs with the right orchestration framework. That's where these platforms come in.

1. LangChain: The Community Favorite

Best for: Developers who need maximum flexibility and community support

If you've been anywhere near AI development lately, you've heard of LangChain. It's become the de facto framework for building LLM applications, and for good reason.

What Makes It Special

LangChain gives you modular components to chain together LLM calls, tools, and external data. Want your agent to search the web, query a database, and then reason about the results? LangChain makes this surprisingly straightforward.

The recent LangGraph extension takes things further by enabling stateful agents with streaming output. This means your agent can maintain complex workflows and provide real-time updates as it thinks through problems.

The Reality Check

Here's what nobody tells you: **LangChain can be resource-intensive**. Complex chains with multiple integrations can rack up API costs fast. You'll also need to stay on top of frequent updates — the ecosystem moves quickly, and keeping integrations in sync takes effort.

One more thing: while basic chains are straightforward, debugging complex agent behaviors requires patience. Expect some trial and error.

Real-World Impact

Klarna used a LangChain-based agent to cut customer support resolution time by 80%. That's not a typo. The framework powers everything from chatGPT-style assistants to sophisticated research tools that combine multiple data sources.

Cost: Free (Apache-2.0), pay only for underlying model API calls

2. OpenAI Agents SDK: Lightweight and Powerful

Best for: Python developers who want elegant multi-agent systems

Released in March 2025, OpenAI's Agents SDK caught everyone by surprise. It's refreshingly minimal — you can spin up multiple collaborating agents in remarkably few lines of code.

The Standout Feature

Despite coming from OpenAI, this SDK is **provider-agnostic**. It orchestrates over 100 different LLMs, not just GPT models. This flexibility is huge for teams that want to experiment with different models or avoid vendor lock-in.

The built-in tracing and debugging tools are exceptional. You can visualize exactly how your agent is thinking through problems, which makes development so much faster.

The Trade-offs

It's new. That means the community is still building out examples and best practices. You might encounter edge cases that aren't well-documented yet.

Also, while it's open-source, some features feel optimized for OpenAI's ecosystem. You'll get the best experience if you're already using OpenAI's platform.

Cost: Free (open-source), pay for model usage

3. Google Agent Development Kit: Low-Code, High-Power

Best for: Google Cloud users building hierarchical agents

Google dropped ADK in April 2025, and it's impressively focused on developer efficiency. The claim? Complex agents in under 100 lines of code. I've tested this, and it's not marketing fluff.

Why It's Different

ADK excels at **hierarchical agents** — essentially, agents made of sub-agents. Picture a supervisor agent that delegates tasks to specialist agents, each handling different aspects of a problem. This mirrors how real teams work, and it's surprisingly effective.

The integration with Google's ecosystem (Gemini, Vertex AI, BigQuery) is seamless. If you're already on Google Cloud, ADK feels like a natural extension of your stack.

The Limitations

Here's the reality: ADK really shines when you're in Google's ecosystem. Outside of it? You'll be building a lot of custom connectors for non-Google services.

The framework is also relatively new with ~10k GitHub stars. While documentation is solid, the community-contributed tools and examples are still growing.

Cost: Free framework, pay for Google Cloud services usage

4. Microsoft AutoGen: Multi-Agent Conversations

Best for: Complex workflows requiring agent collaboration

AutoGen is Microsoft's answer to multi-agent orchestration. Released in late 2023, it's matured significantly and now powers some seriously sophisticated workflows.

The Core Concept

AutoGen uses an event-driven architecture where agents communicate through natural language. One agent might plan a solution while another executes it, with

both iterating until the task is complete.

Think of it like an AI team meeting where different agents play different roles — a Planner, an Executor, a QA specialist — all chatting to solve a problem together.

What You Should Know

AutoGen requires comfort with Python and asynchronous programming. The learning curve is moderate — easier than building from scratch, but steeper than no-code platforms.

Multi-agent systems can be tricky to debug. Agents might miscommunicate or loop endlessly without careful prompt design. AutoGen provides event logs to help, but expect some iteration.

Proven in Production

Pharmaceutical company Novo Nordisk used AutoGen to streamline data science pipelines. The multi-agent approach allowed them to automate complex, regulated workflows that previously required significant manual coordination.

Cost: Free (MIT license), pay for AI model queries

5. Microsoft Semantic Kernel: AI Skills in Your Apps

Best for: Embedding AI features into existing applications

While LangChain and AutoGen focus on building standalone agents, Semantic Kernel takes a different approach: it helps you **add AI capabilities to existing software.**

The Unique Angle

Semantic Kernel treats AI capabilities as "skills" — reusable functions you can compose and call from regular application code. Want to add a "SummarizeEmail" feature to your app? Define it as a skill, and invoke it like any other function.

This bridges traditional software and AI beautifully. Your .NET, Python, or Java application can seamlessly incorporate LLM-powered features without requiring a complete architectural overhaul.

The Trade-offs

The framework has a .NET flavor (it's Microsoft, after all). Python developers might need time to understand patterns like dependency injection.

Also, Semantic Kernel is more about augmenting apps with AI than building fully autonomous agents. For free-roaming agents that handle arbitrary tasks, other frameworks might fit better.

Real-World Applications

Many teams use Semantic Kernel to build Copilot-like assistants within their products. One enterprise software vendor embedded a helpdesk assistant that reads user context and performs actions via APIs — all integrated cleanly into their existing codebase.

Cost: Free, pay for LLM and infrastructure costs

6. Dify: Democratizing AI Agent Development

Best for: Non-developers and rapid prototyping

With over 90,000 GitHub stars by mid-2025, Dify has struck a chord. It's a low-code platform that lets you build AI agents through a visual drag-and-drop interface.

Why It's Exploding in Popularity

Dify makes agent development accessible to non-engineers. You can design workflows visually, connect to hundreds of LLM models, and deploy your agent as

an API or chat widget — all without writing much code.

The platform includes built-in support for RAG (Retrieval-Augmented Generation), ReAct prompting, and function calling. These patterns that usually require careful implementation are just toggle switches in Dify.

The Constraints

While Dify covers common use cases beautifully, highly specialized logic can be challenging to implement purely through the visual interface. You might need custom plugins for edge cases.

Performance can also lag slightly compared to hand-optimized solutions. The abstraction layer trades some efficiency for convenience.

Real-World Success

Startups and non-profits love Dify for quick AI solutions. One e-commerce company built a shopping assistant that answers product questions by pulling from their database — all configured in hours, not weeks.

Cost: Free self-hosted (Apache-2.0), managed cloud has free tier

7. AutoGPT: The Autonomous Pioneer

Best for: Experimental projects and learning about AI agents

AutoGPT deserves credit for igniting the autonomous agent craze in early 2023. It showed the world what happens when you let an AI iterate on its own outputs, decompose goals, and use tools autonomously.

What It Does Differently

AutoGPT breaks down high-level goals automatically. Tell it to “research and write a report on EU economics,” and it generates a plan, gathers data, analyzes topics, drafts sections — all on its own.

The persistent memory means it recalls earlier findings when working on later sub-tasks. It’s like watching an AI work through a complex project step by step.

The Honest Reality

AutoGPT is experimental. It can be brilliant for a while, then get stuck in loops or pursue tangents. Early versions were notorious for consuming massive API tokens while accomplishing little.

You’ll need technical skills to run it — it’s command-line driven and requires managing API keys and Python environments. This isn’t a polished product; it’s a powerful experiment.

Why It Still Matters

AutoGPT remains a landmark project. It’s often cited in research and has inspired countless other agent frameworks. If you want to understand how autonomous agents work under the hood, AutoGPT is an excellent learning tool.

Cost: Free software, pay for API calls (which can add up quickly)

8. Hugging Face Transformers Agents

Best for: ML engineers wanting full model control

Hugging Face's Transformers Agents (evolving into SmolAgents) brings the power of the HuggingFace ecosystem to agent development. The key insight? An agent = LLM + tools.

The Advantage

Being part of Hugging Face means access to thousands of pre-trained models. Need OCR? There's a model. Speech recognition? Already there. The framework lets you compose these specialized models as tools your agent can use.

Remarkably, version 2.0 demonstrated that an open Llama-3–70B model with the right agent strategy can outperform GPT-4-based agents on certain benchmarks. This proves you don't need proprietary APIs for powerful results.

What to Expect

You'll need ML familiarity. The framework assumes you understand model selection, handling outputs, and possibly fine-tuning. It's more code-heavy than no-code platforms.

Running multiple large models (an LLM plus tool models) requires significant compute. Without optimized infrastructure, agents might be slow or expensive.

Where It Shines

The framework excels at multi-modal agents — think an agent that can see (via image models), understand text, and take actions. Perfect for e-commerce product Q&A, healthcare applications, or research assistants.

Cost: Free (Apache license), pay for compute resources

9. Salesforce Agentforce: CRM-Native Intelligence

Best for: Salesforce customers automating customer workflows

Agentforce is Salesforce's 2025 answer to AI agents, deeply integrated into their CRM ecosystem. If you're already running on Salesforce, this is worth serious consideration.

The Integration Story

Agentforce connects natively to Salesforce data — accounts, leads, cases, customer histories. Agents can retrieve and update records, trigger workflows, and incorporate real-time customer context seamlessly.

Pre-built agents handle common scenarios like lead qualification, customer service triage, and marketing campaign suggestions. These serve as templates you can customize through a low-code builder.

The Business Case

For Salesforce shops, Agentforce doesn't require reinventing the wheel. Major clients like Adecco Group, OpenTable, and Saks are already using it, demonstrating real business impact.

The security and governance controls respect Salesforce's role-based access and encryption — crucial for regulated industries.

The Lock-In Factor

Let's be direct: Agentforce requires a Salesforce subscription. It's deeply tied to the Salesforce ecosystem. For non-Salesforce users, it's not relevant.

Pricing likely follows Salesforce's premium model. While exact costs vary, expect enterprise-level pricing that adds to your existing Salesforce spend.

Cost: Part of Salesforce offerings (premium pricing)

10. IBM Watsonx Assistant: Enterprise-Grade Conversations

Best for: Regulated industries requiring strict compliance

IBM's Watsonx Assistant brings decades of AI experience to conversational agents, now infused with modern LLM capabilities. It's built for enterprises that absolutely cannot compromise on security and compliance.

Why Enterprise Trust It

Watsonx Assistant can be hosted on IBM Cloud (including dedicated instances) or on-premises. It offers encryption, access control, and **no data sharing by default** — IBM doesn't train models on your conversations unless you opt in.

The hybrid approach combines rule-based flows with generative AI. Critical transactions like password resets can follow deterministic paths while general queries leverage LLMs.

The Enterprise Tax

IBM solutions are premium priced. Complex setups often require IBM services or skilled internal teams.

The innovation pace feels slower than open-source projects. IBM emphasizes reliability over cutting-edge features, which means new techniques take longer to appear.

Proven Deployment

Banks, insurance companies, and healthcare providers use Watsonx Assistant for customer support. Air Canada's chatbot handles massive query volumes automatically, improving response times while maintaining strict compliance.

Cost: Enterprise pricing (typically usage-based subscriptions)

How to Choose: A Decision Framework

After exploring all these platforms, here's how to think about your choice:

Choose LangChain if: You need maximum flexibility, have a technical team, and want the largest community support.

Choose OpenAI Agents SDK if: You want elegant multi-agent systems with excellent debugging tools and don't mind a newer framework.

Choose Google ADK if: You're on Google Cloud and need hierarchical agents with minimal code.

Choose Microsoft AutoGen if: You need agents that collaborate through conversation to handle complex workflows.

Choose Semantic Kernel if: You want to add AI features to existing applications rather than build standalone agents.

Choose Dify if: You need rapid prototyping or have non-technical team members building agents.

Choose AutoGPT if: You're experimenting, learning, or researching autonomous agent behaviors.

Choose HF Transformers if: You need multi-modal agents with full control over model selection and can handle ML complexity.

Choose Salesforce Agentforce if: You're already on Salesforce and need CRM-integrated agents.

Choose IBM Watsonx if: You're in a regulated industry requiring strict compliance and enterprise support.

The Bottom Line

The AI agent landscape in 2025 is mature, diverse, and incredibly capable. We've moved beyond the experimental phase into production-ready frameworks that real businesses depend on.

Here's what I've learned: **there's no single "best" platform**. Success depends on matching the tool to your needs:

- Startups often thrive with LangChain's flexibility or Dify's speed
- Enterprises gravitate toward vendor-backed solutions like Google ADK, Salesforce, or IBM
- Open-source purists find power in Hugging Face or AutoGPT
- Teams already in ecosystems naturally choose OpenAI, Google, or Microsoft frameworks

The real magic happens when you stop looking for the perfect platform and start building. Pick one that matches your technical capabilities and constraints, then iterate. Every platform on this list can create remarkable AI agents — the differentiator is your implementation.

What are you building? I'd love to hear about your agent projects and which platforms you're exploring. The future is autonomous, and it's being built right now.

. . .

Want to dive deeper into any of these platforms? Check out the official documentation and GitHub repos linked throughout the article. And if you found this guide helpful, share it with your team — choosing the right foundation can save months of work.

Micheal Lanham | Ai Agents In Action | Agent Platform Comparison | Ai Agents | Langchain

Responses (3)

Bgerby

What are your thoughts?

Abyssal Callow H
Oct 23

First rate - and beautiful illustrations!

Reply

Raf Vantongerloo
Oct 9

Great article, fun read!

Reply

Jonathan Matlock
Oct 8

Curious where you see Botpress or Airia fall in line with some of these like Dify.

Reply

More from Micheal Lanham

Micheal Lanham

AI Agents Just Went Enterprise: October 2025 Changed Everything

How Microsoft, OpenAI, Google, and AWS turned “agentic AI” from GitHub experiments into products your company can actually buy

Oct 13 22 2

Micheal Lanham

Building Intelligent RAG Systems with AI Agents: 25 Game-Changing Tips

Your complete guide to creating production-ready Retrieval-Augmented Generation applications that actually work

Sep 24 34

Micheal Lanham

Building Production-Ready RAG Agents: A Developer's Complete Guide

From keyword search to vector embeddings—everything you need to know about implementing Retrieval-Augmented Generation that actually...

Jun 23 28

Micheal Lanham

Building Smarter AI Agents: 7 Game-Changing RAG Strategies That Actually Work

The insider secrets ML engineers use to create AI agents that don't hallucinate and actually know what they're talking about

Sep 14 97 2

See all from Micheal Lanham

Recommended from Medium

In Towards AI by Gao Dalie (高達烈)

RAG is Not Dead! No Chunking, No Vectors, Just Vectorless to Get the Higher Accuracy

Over the past two years, I have written numerous articles on how Retrieval-Augmented Generation has become a standard feature in nearly all...

Oct 17 699 6

In Dare To Be Better by Max Petrusenko

Claude Skills: The $3 Automation Secret That's Making Enterprise Teams Look Like Wizards

How a simple folder is replacing $50K consultants and saving companies literal days of work

Oct 17 398 6

In Artificial Intelligence in Plain English by Alpha Iterations

Build Agentic RAG using LangGraph

A practical guide to build Agentic RAG with complete project code

Oct 19 217 1

In AI Software Engineer by Joe Njenga

This Viral DeepSeek OCR Model Is Changing How LLMs Work

This DeepSeek OCR model hit an overnight success not seen in any other release—4k+ GitHub stars in less than 24 hours and more than 100k...

Oct 23 272 4

Micheal Lanham

AI Agents Just Went Enterprise: October 2025 Changed Everything

How Microsoft, OpenAI, Google, and AWS turned “agentic AI” from GitHub experiments into products your company can actually buy

Oct 13 22 2

In Coding Nexus by Code Coup

Claude Desktop Might Be the Most Useful Free Tool You’ll Install This Year

I didn’t expect much when I first saw the announcement for Claude Desktop. Another AI wrapper, I thought. Maybe with a shiny UI.

6d ago 246 6

See more recommendations