



| Database | Model | Primary Query Language(s) | Strength |
|---------------------------|-------------------------------|---------------------------------------|--|
| Neo4j | Property Graph | Cypher | Developer ergonomics, ecosystem, knowledge graphs. |
| Oracle Graph | Property Graph / RDF in RDBMS | SQL/PGQ, SPARQL | Enterprise RDBMS features + graph. |
| RushDB | Property Graph (zero-config) | JSON search API (Neo4j underpinnings) | Instant graph from JSON/CSV; rapid prototyping. |
| TigerGraph | Distributed Property Graph | GSQL | Massive parallel analytics, real-time recommendations. |
| JanusGraph | Property Graph on backends | Gremlin (TinkerPop) | Portability & backend flexibility. |
| RedisGraph | In-memory Property Graph | Cypher (subset) | Ultra-low latency, Redis integration. |
| Dgraph | Native distributed graph | GraphQL / DQL | GraphQL-first APIs & scale. |
| Neptune, ArangoDB, Nebula | Mixed | SPARQL / Gremlin / AQL / nGQL | Managed AWS service, multi-model, distributed graph. |

Vishal Mysore

Graph Databases & Query Languages in 2025 — A Practical Guide

6 min read · 17 hours ago



Vishal Mysore

Follow

Listen

Share

More

Neo4j, Oracle Graph, RushDB, TigerGraph, and other leading graph systems compared for architects and developers

Applications that need relationship-aware intelligence (recommendations, fraud detection, knowledge graphs, graph-native ML, transit systems) increasingly rely on graph databases. But not all graph stores are the same: they differ in data model, scaling architecture, query language, analytic capabilities, and ecosystem. This article summarizes the state of several major graph databases and the query languages they expose — with practical guidance on when to use which.

• • •

Summary

- Neo4j — Best for developer productivity, property-graph modeling, rich ecosystem and Cypher query language.
- Oracle Graph / Oracle Database (Property Graphs + SQL/PGQ) — Enterprise-grade, brings graph capabilities into a converged RDBMS with SQL/PGQ support for property graphs; great if you need strong transactional RDBMS features plus graph.
- RushDB — New “zero-config” / instant graph DB built on Neo4j that auto-normalizes JSON/CSV to a property graph; targets fast developer onboarding for AI/modern apps. Good for prototypes and teams that want minimal schema work.
- TigerGraph — Designed for large-scale graph analytics; its proprietary GSQL is a high-level, Turing-complete, SQL-like language optimized for parallel graph algorithms. Use it for heavy analytics and production recommendations.
- Gremlin (TinkerPop) / JanusGraph — Gremlin is a graph traversal language supported by JanusGraph and many systems — good when portability across backends is required.
- Dgraph, RedisGraph, Nebula, ArangoDB, Amazon Neptune — each provides distinctive tradeoffs (native GraphQL or DQL, in-memory speed, multi-model, SPARQL/Gremlin support). Pick by API preferences and scale needs.

• • •

What to compare when choosing a graph database

1. Data model: property graph vs RDF (semantic web) vs multi-model
2. Query language: expressiveness, standardization, learning curve (Cypher, GSQL, Gremlin, SPARQL, AQL, DQL)
3. Scale & performance: single-node vs distributed & parallel analytics

4. **Analytics & graph algorithms:** builtin traversal, path-finding, centrality, community detection
5. **Operational needs:** ACID, backups, cloud offering, enterprise support
6. **Ecosystem:** drivers, ORMs, visualization, integrations with ML/LLMs

. . .

Deep dives: systems & languages

Neo4j — Cypher (declarative property-graph)

What it is: Industry-leading property-graph DB focused on developer ergonomics, ACID transactions, and an expressive declarative graph query language called **Cypher**. Neo4j emphasizes pattern matching and readable queries.

Query language: Cypher — declarative, SQL-like for graphs; now part of GQL discussions and supported by openCypher.

Example:

```
MATCH (p:Patient)-[:HAS_DIAGNOSIS]->(d:Disease {name: "Diabetes"})
RETURN p.name, d.name;
```

When to choose Neo4j: Rapid development, tight tooling (Bloom, Aura cloud, drivers for many languages), strong community, and use cases like knowledge graphs, fraud detection, product graphs.

. . .

Oracle Graph / Property Graphs in Oracle Database — SQL/PGQ & converged platform

What it is: Oracle exposes graph capabilities inside Oracle Database (including “property graphs”) and supports **SQL/PGQ** (SQL Property Graph Queries — an ISO/IEC graph query extension) and SPARQL for RDF. This lets you use native SQL tools and enterprise features (security, backup, scale) together with graph analytics.

Query languages supported:

- **SQL/PGQ** – SQL extension for property graph queries (standardizing ways to express graph traversals using SQL)
- **PGQL** on some Oracle tooling / SPARQL for RDF where applicable. Example (conceptual SQL/PGQ):

```
SELECT *
FROM MATCH (p:Person)-[r:FRIEND_OF*1..2]->(q:Person)
WHERE p.id = 'U123';
```

When to choose Oracle Graph: Existing Oracle footprint, need enterprise database features plus graph querying and analytics without deploying a separate graph DB.

• • •

RushDB — Zero-config / Instant Graph (developer-first)

What it is: RushDB is an emergent open-source/ commercial project that promises a zero-config, instant graph database experience: push JSON/CSV, and RushDB auto-normalizes into a labeled meta property graph. It's built with Neo4j under the hood (or integrates with it) and targets fast development for AI applications.

Documentation and vendor pages emphasize no-schema ingestion and a JSON search API.

Query surface: RushDB exposes a JSON-centric query/search API (developer-friendly), rather than forcing learning a new graph DSL. Good for teams that want graph benefits without upfront modeling. Example use-case: rapid prototyping of RAG/LLM knowledge graphs from JSON datasets.

When to choose RushDB: Rapid prototyping, projects where incoming data is heterogeneous JSON and you want immediate graph affordances without schema design.

• • •

TigerGraph — GSQL (parallel analytics)

What it is: TigerGraph is designed for large-scale, distributed graph analytics with built-in parallelism and performance optimizations for multi-hop traversals and graph algorithms. It targets production analytics, real-time recommendations, and ML feature extraction.

Query language: GSQL — SQL-like, Turing-complete language built for expressing iterative graph algorithms, with strong support for bulk-synchronous parallel execution. Example (snippet concept):

```
CREATE QUERY shortestPath(VERTEX<user> start, VERTEX<user> end) {  
    /* GSQL algorithmic logic */  
}
```

When to choose TigerGraph: You need industrial-scale analytics, real-time recommendations, or to run complex graph algorithms in production with low latency.

• • •

JanusGraph / Apache TinkerPop — Gremlin (traversal language)

What it is: JanusGraph is an open-source distributed graph DB that often sits on scalable storage backends (Cassandra, HBase). It implements Apache TinkerPop stack so applications can use Gremlin — a traversal-based, imperative query API that runs across multiple graph engines.

Query language: Gremlin — a path/traversal DSL (imperative style) that's portable across engines that implement TinkerPop.

Example:

```
g.V().hasLabel('person').has('name', 'Alice').out('knows').values('name')
```

When to choose Gremlin/JanusGraph: Need portability across multiple storage backends, or an imperative traversal API for complex traversals and programmatic graph processing.

• • •

RedisGraph — Cypher-like in-memory graph

What it is: RedisGraph embeds graph capabilities into Redis using a compact, adjacency-matrix style representation for high-performance in-memory queries. It supports a subset of Cypher. Good for extremely low-latency lookups and transient graph workloads.

When to choose RedisGraph: Ultra-low latency, ephemeral graphs, or when you already use Redis and need fast graph ops.

• • •

Dgraph — DQL / GraphQL-native

What it is: Dgraph is a distributed native graph database that exposes GraphQL as a first-class API and also its proprietary DQL (previously GraphQL+). It targets developers who prefer GraphQL for graph queries and APIs.

When to choose Dgraph: You want GraphQL-native graph APIs with horizontal scaling and developer-friendly schemas.

• • •

Other notable engines & languages

- **Amazon Neptune:** Managed graph service supporting SPARQL (RDF) and Gremlin (property graph). Good for AWS-native workloads.
- **ArangoDB:** Multi-model DB (document + graph) using AQL (Arango Query Language).
- **NebulaGraph:** Distributed graph DB with nGQL (Cypher-like). [NebulaGraph](#)

• • •

Quick comparison table

| Database | Model | Primary Query Language(s) | Strength |
|---------------------------|-------------------------------|---------------------------------------|--|
| Neo4j | Property Graph | Cypher | Developer ergonomics, ecosystem, knowledge graphs. |
| Oracle Graph | Property Graph / RDF in RDBMS | SQL/PGQ, SPARQL | Enterprise RDBMS features + graph. |
| RushDB | Property Graph (zero-config) | JSON search API (Neo4j underpinnings) | Instant graph from JSON/CSV; rapid prototyping. |
| TigerGraph | Distributed Property Graph | GSQL | Massive parallel analytics, real-time recommendations. |
| JanusGraph | Property Graph on backends | Gremlin (TinkerPop) | Portability & backend flexibility. |
| RedisGraph | In-memory Property Graph | Cypher (subset) | Ultra-low latency, Redis integration. |
| Dgraph | Native distributed graph | GraphQL / DQL | GraphQL-first APIs & scale. |
| Neptune, ArangoDB, Nebula | Mixed | SPARQL / Gremlin / AQL / nGQL | Managed AWS service, multi-model, distributed graph. |

Vishal Mysore

• • •

Practical examples — short snippets

Cypher (Neo4j):

```
CREATE (a:Person {name:'Alice'}),(b:Person {name:'Bob'}),  
      (a)-[:KNOWS]->(b);  
MATCH (a:Person {name:'Alice'})-[:KNOWS]->(b) RETURN b.name;
```

Gremlin (JanusGraph/TinkerPop):

```
g.addV('person').property('name','Alice')  
g.V().has('name','Alice').out('knows').values('name')
```

GSQL (TigerGraph) — conceptual:

```
CREATE QUERY getFriends(VERTEX<person> p) FOR GRAPH MyGraph {  
    start = {p};  
    res = SELECT t FROM start:s - (knows:e) -> :t;  
    PRINT res;  
}
```

SPARQL (RDF example):

```
SELECT ?person WHERE {  
    ?person a ex:Person .  
    ?person ex:knows ex:Bob .  
}
```

DQL / GraphQL (Dgraph):

```
{  
  queryPerson(func: eq(name, "Alice")) {  
    name  
    friend { name }  
  }  
}
```

How query language affects development & portability

- **Declarative (Cypher, SPARQL, AQL):** Easier to read and maintain; good for business queries and analytics.
- **Traversal/Imperative (Gremlin):** More programmatic, flexible for complex traversals and streaming-like graph processing.
- **Turing-complete (GSQL):** Enables writing complex algorithms directly inside the DB — useful for advanced analytics but ties you to the vendor.
- **GraphQL/DQL:** Great for API-first teams that want a single programmable interface.

If portability matters, prefer TinkerPop/Gremlin or standard approaches (SPARQL for RDF). If productivity is key, Cypher and GraphQL approaches win.

• • •

Recommendations by use-case

- Developer prototyping, knowledge graphs, RAG for LLMs: Neo4j (Cypher) or RushDB for zero-config ingestion and fast turn-up.
- Large-scale real-time recommendations / heavy graph ML pipelines: TigerGraph (GSQL) or a distributed system with strong parallelism.
- Enterprise transactional + graph analytics inside existing RDBMS: Oracle Graph (SQL/PGQ) to avoid moving data out of the database.
- Cloud-native, managed option on AWS: Amazon Neptune (SPARQL + Gremlin).
- Graph APIs + front-end teams using GraphQL: Dgraph or ArangoDB.
Hypermode
- Ultra-low-latency in-memory needs: RedisGraph.

• • •

Caveats & vendor lock-in

- Vendor-specific languages (GSQL, nGQL) can offer great capabilities but increase lock-in. If you value portability, prefer open standards (SPARQL, Gremlin, openCypher/GQL-compliant syntax).
- Newer tools like RushDB aim to reduce schema friction — a tradeoff exists between convenience and control/observability.

• • •

Final thoughts — matching language to problem

Pick the graph database based on:

1. Your query patterns (path traversals? analytical algorithms? ad-hoc queries?)

2. Scale & latency needs (single-node vs distributed)
3. Ecosystem fit (cloud provider, existing RDBMS, drivers for your stack)
4. Long-term portability or openness (standards vs proprietary features)

If you're building knowledge graphs to ground LLMs and reduce hallucination, start with Neo4j or RushDB for fast ingestion and Cypher-led querying; for production-grade analytics at scale consider TigerGraph or a distributed graph with Gremlin-compatible tooling for portability.

[Follow](#)

Written by Vishal Mysore

755 followers · 4 following

Holder of multiple patents in AI and software engineering. Passionate about building scalable systems, optimizing performance, & driving AI-powered innovation.

No responses yet



Bgerby

What are your thoughts?

More from Vishal Mysore

 Vishal Mysore

What is BMAD-METHOD™? A Simple Guide to the Future of AI-Driven Development

The BMAD-METHOD™ (Breakthrough Method of Agile AI-Driven Development) framework offers a surprisingly simple, yet highly innovative...

Sep 8  187  2



...

 Vishal Mysore

GitHub Spec Kit vs BMAD-Method: A Comprehensive Comparison : Part 1

GitHub's open-source Spec Kit is a toolkit for spec-driven development that provides a structured process to bring specification-first...

Sep 15  69  4



...

 Vishal Mysore

Comprehensive Guide to Spec-Driven Development Kiro, GitHub Spec Kit, and BMAD-METHOD

What is Spec-Driven Development (SDD)?

Sep 30  16



...



Vishal Mysore

RAG 2.0 : Advanced Chunking Strategies with Examples.

Text chunking is a critical preprocessing step in RAG Based Large Language Model (LLM) applications. While basic strategies like...

Oct 2 24



...

See all from Vishal Mysore

Recommended from Medium



In Medialesson by Marius Schröder

JSON vs TOON—A new era of structured input?

Why structure matters more than ever

Nov 3 163 9



...

 In AI Software Engineer by Joe Njenga

Anthropic Just Solved AI Agent Bloat—150K Tokens Down to 2K (Code Execution With MCP)

Anthropic just released smartest way to build scalable AI agents, cutting token use by 98%, shift from tool calling to MCP code execution

 4d ago  417  33



...

 Vishal Mysore

BMAD-METHOD™ Universal Agent Framework: Expansion Packs That Transform AI Into Any Domain

The BMAD-METHOD™ Universal Agent Framework with Expansion Packs is redefining how AI frameworks scale across industries. Unlike traditional...

Sep 23  8



...

 CodeOrbit

I Built a RAG System for 100,000 Documents—Here's the Architecture

My production system crashed at 2 AM because I underestimated vector databases.

 Nov 1  583  15



...



In Towards AI by Devashish Datt Mamgain

Pseudo-Knowledge Graphs for Better RAG

Retrieval-Augmented Generation (RAG) was supposed to give Large Language Models perfect memory: ask a question, fetch the exact facts, and...

Oct 30 35



...

Arvind Kumar

Building a Ticketing System: Concurrency, Locks, and Race Conditions

What happens when 100,000 fans try to book the same concert ticket at exactly 10:00 AM? Let's design a ticketing system that prevents...

Oct 30 616 8



...

[See more recommendations](#)