

Szolkowski Blog

Blog with Episerver/Optimizely and tech stuff

[Tags](#)
[About](#)
[Archive](#)

8 October 2025

/ Quiet Performance Wins: Scheduled Job for SQL Index Maintenance in Optimizely

by Stanisław Szolkowski

Tags: [database \(1\)](#) [db \(1\)](#) [episerver \(8\)](#) [indexes \(1\)](#)
[jobs \(3\)](#) [maintenance \(2\)](#) [optimizely \(8\)](#)
[performance \(1\)](#) [scheduled jobs \(2\)](#)

As Optimizely CMS projects grow, it's not uncommon to introduce custom tables—whether for integrations, caching, or specialized business logic. But with great schema comes great responsibility: SQL Server indexes and statistics need love too.

While Optimizely handles its own data structures well, custom tables can quietly degrade performance if left unchecked. Optimizely Commerce includes a built-in job for index maintenance, but if your solution uses only CMS, this functionality will be missing. In this post, I'll show how to automate index and statistics maintenance using a scheduled job.

// Why You Should Care

SQL Server relies heavily on up-to-date statistics and healthy indexes to optimize query execution. Fragmented indexes and stale stats can lead to slow queries, increased CPU usage, and unhappy editors.

If you're adding custom tables to your CMS database, especially ones that grow over time, you should consider regular maintenance. And what better way than a scheduled job that runs quietly in the background?

// The Scheduled Job

Here's a simple implementation of a scheduled job that performs index and statistics maintenance on selected custom tables. You can trigger it manually or schedule it via Optimizely's job system.

[View Fullscreen

Copy

```
/// <summary>
/// Automated database index maintenance job that runs o
```

```
/// This job analyzes index fragmentation and performs maintenance
/// </summary>
[ScheduledPlugIn(
    DisplayName = "Database Index Maintenance Scheduled Job",
    SortIndex = 20000)]
public sealed class DatabaseIndexMaintenanceScheduledJob
{
    private bool _stopRequested;
    private readonly IConfiguration _configuration;

    /// <summary>
    /// Constructor injecting configuration for database connection
    /// Sets IsStoppable to allow manual termination of the job
    /// </summary>
    public DatabaseIndexMaintenanceScheduledJob(IConfiguration configuration)
    {
        _configuration = configuration;

        // Allow administrators to stop the job if it's running
        IsStoppable = true;
    }

    /// <summary>
    /// Handles stop requests by setting a flag that's checked in ExecuteInternal
    /// This allows graceful cancellation between maintenance steps
    /// </summary>
    public override void Stop()
    {
        _stopRequested = true;
    }

    /// <summary>
    /// Main entry point for the scheduled job execution
    /// Retrieves the database connection string and performs maintenance
    /// </summary>
    public override string Execute()
    {
        // Get the connection string from configuration
        var connectionString = _configuration.GetConnectionString("Default");

        // Validate connection string exists before proceeding
        var result = !string.IsNullOrEmpty(connectionString)
            ? ExecuteInternal(connectionString)
            : "Connection string is empty";

        return result;
    }

    /// <summary>
    /// Core maintenance logic that analyzes and optimizes indexes
    /// Uses a three-phase approach:
    /// 1. Query all indexes and measure their fragmentation
    /// 2. Rebuild or reorganize indexes based on fragmentation
    /// 3. Update statistics for tables with fragmented indexes
    /// </summary>
    private string ExecuteInternal(string connectionString)
    {
        // StringBuilder accumulates log messages for the job
        var log = new StringBuilder();
        try
        {
            // Establish database connection using 'using' statement
            using var conn = new SqlConnection(connectionString);
            conn.Open();

            log.AppendLine("Starting index maintenance...");

            // Query SQL Server's Dynamic Management View
            // sys.dm_db_index_physical_stats provides fragmentation information
            var indexQuery = @"
                SELECT ...
                FROM sys.dm_db_index_physical_stats ...
            ";
            ...
        }
        catch (Exception ex)
        {
            log.AppendLine($"Error: {ex.Message}");
        }
    }
}
```

```

SELECT OBJECT_SCHEMA_NAME(s.[object_id])
      ,OBJECT_NAME(s.[object_id]) AS Table
      ,i.name AS IndexName,
      ,s.avg_fragmentation_in_percent AS AvgFrag
  FROM sys.dm_db_index_physical_stats(DB_ID(), object_id, 0, 1, 1)
 JOIN sys.indexes i
    ON s.[object_id] = i.[object_id]
   AND s.index_id = i.index_id
 WHERE i.type_desc IN ('CLUSTERED', 'NONCLUSTERED')
   AND s.page_count > 100;" // Only analyze tables with more than 100 pages

using var cmd = new SqlCommand(indexQuery, connection);
using var reader = cmd.ExecuteReader();

// Phase 1: Collect all indexes and their fragmentation levels
// Store in a list to avoid maintaining an ordered collection
var indexList = new List<(string Schema, string Table, string IndexName, double AvgFrag)>();
while (reader.Read() && !_stopRequested)
{
    var schema = reader.GetString(0);           // Schema
    var table = reader.GetString(1);            // Table
    var index = reader.GetString(2);            // Index Name
    var frag = reader.GetDouble(3);             // Avg Fragmentation

    indexList.Add((schema, table, index, frag));
}

// Close the reader before executing maintenance commands
reader.Close();

// Phase 2: Perform index maintenance based on fragmentation level
// Industry best practices: REBUILD > 30%, REORG > 5%
foreach (var (schema, table, index, frag) in indexList)
{
    // Check for stop request between each iteration
    if (_stopRequested) break;

    // Use pattern matching to determine the appropriate maintenance command
    var sql = frag switch
    {
        // Severe fragmentation (>30%): REBUILD
        // Add optionally "WITH ONLINE = ON"
        > 30 => $"ALTER INDEX [{index}] ON [{table}] REBUILD WITH ONLINE = ON";
        // Moderate fragmentation (5-30%): REORG
        // This is always an online operation
        > 5 => $"ALTER INDEX [{index}] ON [{table}] REORG";
        // Low fragmentation (<5%): No action required
        _ => null
    };

    // Execute the maintenance command if an SQL string was generated
    if (sql != null)
    {
        log.AppendLine($"Maintaining index [{index}] on table [{table}]");
        using var alterCmd = new SqlCommand(sql, connection);
        // Set a generous timeout for long-running operations
        alterCmd.CommandTimeout = 180;
        alterCmd.ExecuteNonQuery();
    }
}

// Phase 3: Update statistics for tables that were maintained
// Statistics help the query optimizer make better decisions
foreach (var (schema, table, _, frag) in indexList)
{
    // Check for stop request between each statistics update
    if (_stopRequested) break;

    var statsQuery = $"UPDATE STATISTICS {schema}.{table} WITH COLUMNS";
    using var statsCmd = new SqlCommand(statsQuery, connection);
    statsCmd.ExecuteNonQuery();
}

```

```

        // Only update statistics for tables with
        // SAMPLE 50 PERCENT balances accuracy with
        var sql = frag switch
        {
            > 5 => $"UPDATE STATISTICS [{schema}].{table} SAMPLE 50 PERCENT";
            _ => null
        };

        if (sql != null)
        {
            log.AppendLine($"Updating statistics for {table}");
            using var statsCmd = new SqlCommand(sql);
            // Set a generous timeout for long-running operations
            statsCmd.CommandTimeout = 180;
            statsCmd.ExecuteNonQuery();
            log.AppendLine("Statistics updated.");
        }
    }

    conn.Close();
}
catch (Exception ex)
{
    // Log any errors that occur during maintenance
    log.AppendLine($"Error: {ex.Message}");
}

// Return the accumulated log as the job execution log
return log.ToString();
}
}

```

// Performance Considerations

/// During Execution

REBUILD operations:

- High CPU usage (50-80% spike for 1-5 minutes)
- Locks the table (unless `ONLINE = ON` on Enterprise Edition)
- Should be run during maintenance windows

REORGANIZE operations:

- Minimal CPU impact (10-20%)
- Online operation (no blocking)
- Safe to run during business hours

/// After Execution

Typical improvements on databases with 30%+ fragmentation:

- Query performance: 15-40% faster
- CPU usage: 10-15% reduction
- Page I/O: 20-30% reduction

Note: Benefits are most noticeable on:

- Tables with > 1M rows
- Queries with table/index scans
- Reports and analytics queries

// What Gets Maintained?

This job analyzes **all indexes** across your entire database, including:

Table Type	Examples	Maintained?
Your custom tables	<code>CustomOrderCache</code> , <code>IntegrationLog</code>	Yes
Optimizely CMS core	<code>tblContent</code> , <code>tblContentProperty</code> , <code>tblWorkContent</code>	Yes
Commerce tables	<code>OrderGroup</code> , <code>Shipment</code> , <code>LineItem</code>	Yes (if installed)
ASP.NET Identity	<code>AspNetUsers</code> , <code>AspNetRoles</code>	Yes

// Is This Safe?

Yes, generally. Index maintenance operations are safe on all tables. However:

- **REBUILD operations** may briefly lock tables on Standard Edition
- Large Optimizely tables (like `tblContentProperty`) may take several minutes
- First run might take 10-20 minutes on established sites

// Should You Filter?

For production safety, consider filtering to custom tables only if:

- You have a very large CMS database (50GB+)
- You're on SQL Server Standard Edition (no ONLINE rebuilds)
- You want to minimize maintenance window impact

// Understanding Statistics Updates

`UPDATE STATISTICS` ensures the query optimizer has accurate data about:

- Row counts
- Data distribution
- Index selectivity

The `SAMPLE 50 PERCENT` option:

- Faster than FULLSCAN
- Accurate enough for most scenarios
- Use `WITH FULLSCAN` for critical tables if needed

// Notes

`UPDATE STATISTICS` ensures the query planner has fresh data to work with.

You can extend the job to log execution time or errors using Optimizely's logging framework.

Provided job targets ALL indexes in the database, including Optimizely's core tables. Logic can be adjusted to target only whitelisted tables by changing index query:

```
SELECT OBJECT_SCHEMA_NAME(s.[object_id]) AS SchemaName,
       OBJECT_NAME(s.[object_id]) AS TableName,
       i.name AS IndexName,
       s.avg_fragmentation_in_percent AS Frag
  FROM sys.dm_db_index_physical_stats(DB_ID(), NULL, NULL)
 JOIN sys.indexes i ON s.[object_id] = i.[object_id]
 WHERE i.type_desc IN ('CLUSTERED', 'NONCLUSTERED')
```

```
AND s.page_count > 100  
AND OBJECT_NAME(s.[object_id]) IN (XXXX)
```

Query can be also adjusted to filter by schema or prefix.

// Troubleshooting

Job times out:

- Increase the SQL command timeout
- Consider running REORGANIZE on specific indexes rather than ALL

Permission errors:

- Ensure the app pool identity has `db_ddladmin` role
- Check Azure SQL firewall rules if using DXP

High CPU during execution:

- Move to off-peak hours
- Add `WITH (ONLINE = ON)` option if using Enterprise edition

// Summary

This kind of job is especially useful in environments where custom tables are updated frequently but not covered by Optimizely's internal maintenance routines. It's a small addition that can yield big performance wins.

If you're deploying to DXP, make sure the job is safe to run in production and doesn't interfere with other scheduled tasks. It is usually good to run this job weekly during the low traffic hours.



Tags

[apple silicon \(2\)](#)
[application insights \(1\)](#)
[arm \(2\)](#)
[background jobs \(1\) ci \(2\)](#)
[database \(1\) db \(1\)](#)
[devops \(2\) dxp \(1\)](#)
[episerver \(8\)](#)
[github \(2\) hangfire \(1\)](#)
[indexes \(1\) jobs \(3\)](#)
[m1 \(2\) maintenance \(2\)](#)
[ms sql server \(1\)](#)
[optimizely \(8\)](#)
[performance \(1\)](#)
[pipeline \(2\)](#)
[scheduled jobs \(2\)](#)
[sonarcloud \(1\)](#)
[sonarqube \(1\)](#)
[sql server \(1\) workflow \(2\)](#)

[Read more articles](#) | subscribe [via RSS](#)